

Supporting Materials for “Automatic Differentiation to Improve Parameter Estimation for Partially Observed Markov Processes”

Kevin Tan¹, Aaron J. Abkemeier², Giles Hooker¹ and Edward L. Ionides²

¹University of Pennsylvania and ²University of Michigan

Supplementary Content

S1 Proof of Theorem 1 (DMOP-α Functional Forms)	S-2
S2 Proof of Theorem 2 (Optimization Convergence Analysis)	S-5
S3 Feynman-Kac Models and Monte Carlo Approximations	S-10
S4 Proof of Theorem 3 via a Strong Law of Large Numbers for Triangular Arrays of Particles With Off-Parameter Resampling	S-12
S5 Proof of Theorem 4 via Consistency of Off-Parameter Resampled Gradient Estimates	S-18
S6 Proof of Theorem 5 (Bias-Variance Analysis)	S-19
S7 Sharper Bias-Variance Analysis	S-29
S8 Figures for Bayesian Inference	S-50
S9 Justification for the Choice of Baseline Comparisons	S-53
S10 Algorithmic Parameters	S-54

S1 Proof of Theorem 1 (DMOP- α Functional Forms)

Theorem 1 follows immediately as a consequence of the following results, Lemmas S1 and S2. Here, an A superscript is used to denote ancestral trajectories, so that $x_{1:n,j}^{A,P,\theta}$ contains the history of the prediction particle $x_{n,j}^{P,\theta}$ at observation times $1:n$. Similarly, $g_{i,j}^{A,P,\theta}$ is the ancestral weight of $x_{n,j}^{P,\theta}$ at observation time i .

Lemma S1. Write $\nabla_{\theta} \hat{\ell}^{\alpha}(\theta)$ for the DMOP- α gradient estimate (or, equivalently, the gradient of MOP- α when $\theta = \phi$). Consider the case where we use the after-resampling conditional likelihood estimate so that $\hat{\mathcal{L}}(\theta) = \prod_{n=1}^N L_n^{A,\theta,\alpha}$. When $\alpha = 1$,

$$\nabla_{\theta} \hat{\ell}^1(\theta) = \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \log f_{Y_{1:N}|X_{1:N}}(y_{1:N}^* | x_{1:n,j}^{A,F,\theta}; \theta), \quad (\text{S1})$$

yielding the bootstrap filter estimator of Poyiadjis et al. (2011) and Ścibior and Wood (2021).

Proof. Consider the case of MOP- α when $\alpha = 1$ and $\theta = \phi$. We then have a nice telescoping product property for the after-resampling likelihood estimate:

$$\begin{aligned} \hat{\mathcal{L}}^1(\theta) &:= \prod_{n=1}^N L_n^{A,\theta,\alpha} = \prod_{n=1}^N L_n^{\phi} \cdot \frac{\sum_{j=1}^J w_{n,j}^{F,\theta}}{\sum_{j=1}^J w_{n,j}^{P,\theta}} = \prod_{n=1}^N L_n^{\phi} \cdot \frac{\sum_{j=1}^J w_{n,j}^{F,\theta}}{\sum_{j=1}^J w_{n-1,j}^{F,\theta}} \\ &= \left(\frac{1}{J} \sum_{j=1}^J w_{N,j}^{F,\theta} \right) \prod_{n=1}^N L_n^{\phi}, \end{aligned} \quad (\text{S2})$$

where the third equality follows from the choice of $\alpha = 1$, and the fourth equality is the resulting telescoping property. The log-derivative identity lets us decompose the score estimate as

$$\nabla_{\theta} \hat{\ell}^1(\theta) = \frac{\nabla_{\theta} \hat{\mathcal{L}}^1(\theta)}{\hat{\mathcal{L}}^1(\theta)} = \frac{\nabla_{\theta} \left(\frac{1}{J} \sum_{j=1}^J w_{N,j}^{F,\theta} \right) \prod_{n=1}^N L_n^{\phi}}{\prod_{n=1}^N L_n^{\phi}} = \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} w_{N,j}^{F,\theta}. \quad (\text{S3})$$

From (S3), we see that the derivative of the log-likelihood estimate is

$$\nabla_{\theta} \hat{\ell}^1(\theta) := \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} w_{N,j}^{F,\theta}. \quad (\text{S4})$$

We proceed to decompose (S4). First, observe that as $\alpha = 1$,

$$w_{n,j}^{P,\theta} = w_{n-1,j}^{F,\theta} \frac{g_{n,j}^{\theta}}{g_{n,j}^{\phi}} = \prod_{i=1}^n \frac{g_{i,j}^{A,P,\theta}}{g_{i,j}^{A,P,\phi}}, \quad (\text{S5})$$

Note that the right hand side of (S5) is the cumulative product of measurement density ratios over the ancestral trajectory of the j -th prediction particle at timestep n . We then use the log-derivative identity again, yielding the following expression for the gradient of the log-weights as the sum of the log measurement densities over the ancestral trajectory,

$$\frac{\nabla_{\theta} w_{n,j}^{P,\theta}}{w_{n,j}^{P,\theta}} = \nabla_{\theta} \log w_{n,j}^{P,\theta} = \nabla_{\theta} \log \left(\prod_{i=1}^n \frac{g_{i,j}^{A,P,\theta}}{g_{i,j}^{A,P,\phi}} \right) \quad (\text{S6})$$

$$= \nabla_{\theta} \sum_{i=1}^n \left(\log g_{i,j}^{A,P,\theta} - \log g_{i,j}^{A,P,\phi} \right) \quad (\text{S7})$$

$$= \sum_{i=1}^n \nabla_{\theta} \log g_{i,j}^{A,P,\theta}. \quad (\text{S8})$$

This is equal to the gradient of the logarithm of the conditional density of the observed measurements given the ancestral trajectory of the j -th prediction particle up to timestep n ,

$$\nabla_{\theta} \sum_{n=1}^N \log g_{n,j}^{A,P,\theta} = \nabla_{\theta} \log \left(\prod_{n=1}^N g_{n,j}^{A,P,\theta} \right) \quad (\text{S9})$$

$$= \nabla_{\theta} \log \left(\prod_{n=1}^N f_{Y_n|X_n}(y_n^* | x_{n,j}^{A,P,\theta}; \theta) \right) \quad (\text{S10})$$

$$= \nabla_{\theta} \log f_{Y_{1:N}|X_{1:N}}(y_{1:N}^* | x_{1:n,j}^{A,P,\theta}; \theta). \quad (\text{S11})$$

Multiplying both sides of the expression by $w_{N,j}^{P,\theta}$ yields an expression for the gradient of the weights at timestep N ,

$$\nabla_{\theta} w_{N,j}^{P,\theta} = w_{N,j}^{P,\theta} \sum_{n=1}^N \nabla_{\theta} \log g_{n,j}^{A,P,\theta} = w_{N,j}^{P,\theta} \nabla_{\theta} \log f_{Y_{1:N}|X_{1:N}}(y_{1:N}^* | x_{1:n,j}^{A,P,\theta}; \theta). \quad (\text{S12})$$

Substituting the above identity into the log-likelihood decomposition obtained earlier in Equation S3 yields

$$\begin{aligned} \nabla_{\theta} \hat{\ell}^1(\theta) &= \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} w_{N,j}^{F,\theta} = \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} w_{N,k_j}^{P,\theta} \\ &= \frac{1}{J} \sum_{j=1}^J w_{N,k_j}^{P,\theta} \nabla_{\theta} \log f_{Y_{1:N}|X_{1:N}}(y_{1:N}^* | x_{1:n,k_j}^{A,P,\theta}; \theta). \end{aligned} \quad (\text{S13})$$

Finally, observing that $\theta = \phi$ implies $w_{N,j}^{F,\theta} = 1$, we obtain

$$\nabla_{\theta} \hat{\ell}^1(\theta) = \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \log f_{Y_{1:N}|X_{1:N}}(y_{1:N}^* | x_{1:n,j}^{A,F,\theta}; \theta). \quad (\text{S14})$$

This has the same undiscounted path-space score target as the estimators of Poyiadjis et al. (2011); Ścibior and Wood (2021) for the bootstrap filter. The reparameterization score contribution need not equal their complete-data likelihood-ratio score contribution particle by particle. $\square$

Note that the variance of the DMOP- α log-likelihood derivative estimate scales poorly with N when $\theta \neq \phi$. This can be seen by observing that

$$\nabla_{\theta} \hat{\ell}^1(\theta) = \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} w_{N,j}^{F,\theta} = \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} w_{N,k_j}^{P,\theta}$$

$$= \frac{1}{J} \sum_{j=1}^J w_{N,k_j}^{P,\theta} \nabla_{\theta} \log f_{Y_{1:N}|X_{1:N}}(y_{1:N}^* | x_{1:n,k_j}^{A,P,\theta}; \theta). \quad (\text{S15})$$

When $\theta \neq \phi$, we see that $w_{N,k_j}^{P,\theta} = O(c^N)$. When $\theta = \phi$, a generic normalized path-functional inequality gives the coarse upper bound $O(N^4/J)$ after rescaling to the full unnormalized score. Section S7 instead proves the sharper current-particle upper bound $O(N^3/J)$ at $\alpha = 1$.

Lemma S2. Write $\nabla_{\theta} \hat{\ell}^{\alpha}(\theta)$ for the DMOP- α gradient estimate (or, equivalently, the gradient of MOP- α when $\theta = \phi$). Consider the case where we use the after-resampling conditional likelihood estimate so that $\hat{\mathcal{L}}(\theta) = \prod_{n=1}^N L_n^{A,\theta,\alpha}$. When $\alpha = 0$,

$$\nabla_{\theta} \hat{\ell}^0(\theta) = \frac{1}{J} \sum_{n=1}^N \sum_{j=1}^J \nabla_{\theta} \log \left(f_{Y_n|X_n}(y_n^* | x_{n,j}^{F,\theta}; \theta) \right), \quad (\text{S16})$$

yielding the estimate of Naesseth et al. (2018) when applied to the bootstrap filter.

Proof. First, write

$$s_{n,j} = \frac{f_{Y_n|X_n}(y_n^* | x_{n,j}^{P,\theta}; \theta)}{f_{Y_n|X_n}(y_n^* | x_{n,j}^{P,\phi}; \theta)} \quad (\text{S17})$$

as shorthand for the measurement density ratios. Observe that, when $\alpha = 0$, the likelihood estimate becomes

$$\hat{\mathcal{L}}^0(\theta) := \prod_{n=1}^N L_n^{A,\theta,\alpha} = \prod_{n=1}^N L_n^{\phi} \cdot \frac{\sum_{j=1}^J w_{n,j}^{F,\theta}}{\sum_{j=1}^J w_{n,j}^{P,\theta}} \quad (\text{S18})$$

$$= \prod_{n=1}^N L_n^{\phi} \cdot \frac{1}{J} \sum_{j=1}^J s_{n,j} \quad (\text{S19})$$

$$= \prod_{n=1}^N L_n^{\phi} \cdot \frac{1}{J} \sum_{j=1}^J \frac{f_{Y_n|X_n}(y_n^* | x_{n,j}^{P,\theta}; \theta)}{f_{Y_n|X_n}(y_n^* | x_{n,j}^{P,\phi}; \theta)}. \quad (\text{S20})$$

We lose the nice telescoping property observed in the MOP-1 case, but this expression still yields something useful. This is because its gradient when $\theta = \phi$ is therefore

$$\nabla_{\theta} \hat{\ell}^0(\theta) := \sum_{n=1}^N \nabla_{\theta} \log \left(L_n^{\phi} \frac{1}{J} \sum_{j=1}^J s_{n,j} \right) \quad (\text{S21})$$

$$= \sum_{n=1}^N \frac{\nabla_{\theta} \left(L_n^{\phi} \frac{1}{J} \sum_{j=1}^J s_{n,j} \right)}{\left(L_n^{\phi} \frac{1}{J} \sum_{j=1}^J s_{n,j} \right)} \quad (\text{S22})$$

$$= \sum_{n=1}^N \frac{\sum_{j=1}^J \nabla_{\theta} s_{n,j}}{\sum_{j=1}^J s_{n,j}} \quad (\text{S23})$$

$$= \sum_{n=1}^N \frac{1}{J} \sum_{j=1}^J \frac{\nabla_{\theta} f_{Y_n|X_n}(y_n^* | x_{n,j}^{F,\theta}; \theta)}{f_{Y_n|X_n}(y_n^* | x_{n,j}^{F,\phi}; \phi)} \quad (\text{S24})$$

$$= \frac{1}{J} \sum_{n=1}^N \sum_{j=1}^J \nabla_{\theta} \log \left(f_{Y_n|X_n}(y_n^*|x_{n,j}^{F,\theta}; \theta) \right), \quad (\text{S25})$$

where we use the log-derivative trick in the second equality, observe that $\sum_{j=1}^J s_{n,j} = J$ when $\theta = \phi$ in the fourth equality, and use the log-derivative trick again while noting that $\theta = \phi$ in the fifth equality. This yields the desired result. $\square$

S2 Proof of Theorem 2 (Optimization Convergence Analysis)

The analysis in this section is similar to the approach of Roosta-Khorasani and Mahoney (2016), with the caveat that none of the matrix concentration bounds they use apply here as the particles are dependent. We instead use the concentration inequality from Del Moral and Rio (2011) to bound the gradient and Hessian estimates. In this section, we fix $\omega \in \Omega$ only within each filtering iteration, and analyze Algorithm 3 post-iterated filtering.

The convergence analysis in Theorem 2 is limited to the case where $-\ell$ is γ -strongly convex. Under regularity conditions, we expect the log-likelihood to be approximately quadratic in a neighborhood of its maximum, though in practice the likelihood surfaces for POMP models are often highly nonconvex globally. The convergence to an optimum, local or global, must therefore be sensitive to initialization.

For simplicity, we only provide the argument for DMOP-1 and note that the general case follows as the required ϵ -approximation can still be obtained for sufficiently large α . When $\alpha = 1$, the required J is given in the below lemmas. We start by bounding the error on the gradient estimate, in Lemma S3.

Lemma S3 (Concentration of Measure for Gradient Estimate). *Consider the gradient estimate obtained by MOP-1, which we know by Theorem 4 is consistent for the score, where $\theta = \phi$. For $\|\nabla_{\theta} \hat{\ell}^1(\theta) - \nabla_{\theta} \ell(\theta)\|_2$ to be bounded by ϵ with probability $1 - \delta$, we require*

$$J > \max \left\{ 2G(\theta) \frac{r_N \sqrt{p}}{\epsilon} \left(1 + h^{-1} \left(\log \left(\frac{2p}{\delta} \right) \right) \right), 8G(\theta)^2 \beta_N^2 \frac{p \log(2p/\delta)}{\epsilon^2} \right\}, \quad (\text{S26})$$

where $NG'(\theta) \leq G(\theta)$ are defined in Assumptions (A3) and (A4), $h(t) = \frac{1}{2}(t - \log(1+t))$, and β_N and r_N are two additional finite model-specific constants that do not depend on J , but do depend on N and p , as defined in Del Moral and Rio (2011). Equivalently, with probability at least $1 - \delta$, it holds that

$$\|\nabla_{\theta} \hat{\ell}^1(\theta) - \nabla_{\theta} \ell(\theta)\|_2 \leq G(\theta) \left(\frac{r_N \sqrt{p}}{J} (1 + h^{-1}(\log(2p/\delta))) + \sqrt{\frac{2p \log(2p/\delta)}{J}} \beta_N \right). \quad (\text{S27})$$

According to Del Moral and Rio (2011), under suitable regularity conditions, r_N and β_N are linear in N . The generic path-space result of Del Moral and Doucet (2003) then gives the coarse error bound $O(N^2 J^{-1/2})$, or MSE $O(N^4/J)$, after rescaling to the unnormalized score. This is not the sharp variance bound used in Theorem S5: Section S7 gives $O(N^3/J)$ at $\alpha = 1$ and $O\{N/[J(1-\alpha)^2]\}$ for fixed $\alpha < 1$, up to the common parameter and mixing-time factors.

Proof. We use the concentration inequality of Del Moral and Rio (2011) to bound the deviation of the gradient estimate from the gradient of the negative log-likelihood in the sup norm with a union

bound. Fix $\theta = \phi$. From the decomposition in the proof of Lemma S1, as $w_{N,j}^{F,\theta} = 1$ when $\theta = \phi$, we have that

$$\nabla_{\theta} \hat{\ell}(\theta) := \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} w_{N,j}^{F,\theta} = \frac{1}{J} \sum_{j=1}^J \sum_{n=1}^N w_{N,k_j}^{P,\theta} \nabla_{\theta} \log g_{n,k_j}^{A,\theta} = \frac{1}{J} \sum_{j=1}^J \sum_{n=1}^N \nabla_{\theta} \log g_{n,k_j}^{A,\theta}. \quad (\text{S28})$$

Define $\varphi_n^i(x_{n,j}^{F,\theta}) := \frac{\partial}{\partial \theta_i} \log g_{n,k_j}^{A,\theta}$, which is a functional of the filtering particles $x_{n,j}^{F,\theta} = x_{n,k_j}^{P,\theta}$. These are bounded measurable functionals bounded by $G'(\theta)$ by Assumption (A3). Therefore, these have bounded oscillation, satisfying the requirement that $\text{osc}\left(\frac{\partial}{\partial \theta_i} \varphi_i(x_{n,j}^{P,\theta})\right) \leq G'(\theta)$. Note that Del Moral and Rio (2011) in fact assume $\text{osc}(f) \leq 1$, so we simply scale their bound accordingly.

Now we apply the Hoeffding-type concentration inequality in Theorem 1.2 of Del Moral and Rio (2011), and a union bound over each $\varphi_n^i(x_{n,j}^{F,\theta})$, totaling N timesteps and p parameters, to find that

$$\max_{n=1,\dots,N} \left\| \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \log g_{n,k_j}^{A,\theta} - \nabla_{\theta} \ell_n(\theta) \right\|_{\infty} \leq G'(\theta) \left(\frac{r_N}{J} (1 + h^{-1}(t)) + \sqrt{\frac{2t}{J}} \beta_N \right) \quad (\text{S29})$$

with probability at least $1 - 2Np \exp(-t)$. The concentration inequality justifying (S29) directly concerns the error from the expectation under the filtering distribution. However, we replace this with $\nabla_{\theta} \ell_n(\theta)$ in (S29) since our regularity conditions ensure that the gradient of the measurement density is bounded and thus, by dominated convergence, the expectation of this gradient equals the gradient of its expectation. Summing over N , and noting that $NG'(\theta) \leq G(\theta)$, we therefore obtain, with the same probability, that

$$\left\| \frac{1}{J} \sum_{j=1}^J \sum_{n=1}^N \nabla_{\theta} \log g_{n,k_j}^{A,\theta} - \nabla_{\theta} \ell(\theta) \right\|_{\infty} \leq G'(\theta) N \left(\frac{r_N}{J} (1 + h^{-1}(t)) + \sqrt{\frac{2t}{J}} \beta_N \right) \quad (\text{S30})$$

$$\leq G(\theta) \left(\frac{r_N}{J} (1 + h^{-1}(t)) + \sqrt{\frac{2t}{J}} \beta_N \right). \quad (\text{S31})$$

We split the δ failure probability among these $2Np$ terms, to find $\delta \leq 2Np \exp(-t)$, and therefore, $t \leq \log(2Np/\delta)$, where $h(t) = \frac{1}{2}(t - \log(1+t))$. The two additional model-specific parameters are β_t and r_t , which do not depend on J . The analogous bound for the 2-norm follows from scaling the right-hand side by $\sqrt{p}$, to require

$$\|\nabla_{\theta} \hat{\ell}(\theta) - \nabla_{\theta} \ell(\theta)\|_2 \leq G(\theta) \left(\frac{r_N \sqrt{p}}{J} (1 + h^{-1}(\log(2p/\delta))) + \sqrt{\frac{2p \log(2p/\delta)}{J}} \beta_N \right). \quad (\text{S32})$$

We therefore need

$$J > \max \left\{ 2G(\theta) \frac{r_N \sqrt{p}}{\epsilon} \left(1 + h^{-1} \left(\log \left(\frac{2p}{\delta} \right) \right) \right), 8G(\theta)^2 \beta_N^2 \frac{p \log(2p/\delta)}{\epsilon^2} \right\}. \quad (\text{S33})$$

□

We proceed to study the particle Hessian estimate, providing a guarantee for its positive-definiteness in Lemma S4. This bound is relevant if one chooses to use a second-order method involving a particle Hessian estimate.

Lemma S4 (Minimum Eigenvalue Bound for Hessian Estimate). *Assume that the Hessian of the negative log-likelihood $H = \sum_{j=1}^J \mathbb{E}H_j$ has a minimum eigenvalue $0 < \gamma < 1$, and that $\mathbb{E}\lambda_{\min}(H_j) = \gamma' > 0$. If*

$$J > \max \left\{ \frac{2r_t(1 + h^{-1}(t)) + 2c}{\gamma'}, \frac{2(2t\beta_t^2 + c)^2}{\gamma'^2} \right\} \geq \frac{r_t(1 + h^{-1}(t))}{\gamma'} + \sqrt{2tJ}\beta_t/\gamma' + c/\gamma' \quad (\text{S34})$$

then $\hat{H}(\theta)$ is invertible and positive definite with minimum eigenvalue greater than or equal to $c \in (0, \sum_{j=1}^J \mathbb{E}\lambda_{\min}(H_j))$, with probability at least $1 - \exp(-t)$.

Proof. Write $\hat{H}(\theta) = \hat{H} = \sum_{j=1}^J H_j$ for the estimate of the negative of the Hessian, where H_j is an element of the outer sum over the J particles.

As the negative log-likelihood is convex, we want to bound the minimum eigenvalue of $\hat{H}(\theta)$ from below with high probability, so that all the eigenvalues of $\hat{H}(\theta)$ are positive with high probability. This ensures that the estimated Hessian is invertible and positive-definite.

It is known that the minimum eigenvalue of a symmetric matrix is concave. Therefore, it suffices to show that the first inequality in the below expression

$$0 < \sum_{j=1}^J \lambda_{\min}(H_j) \leq \lambda_{\min} \left(\sum_{j=1}^J H_j \right) = \lambda_{\min}(\hat{H}) \quad (\text{S35})$$

holds with high probability. We apply the particle Hoeffding concentration inequality from Del Moral and Rio (2011) to find that

$$\frac{1}{J} \sum_{j=1}^J \lambda_{\min}(H_j) - \mathbb{E}_{\tilde{\pi}_t} \lambda_{\min}(H_j) = \frac{1}{J} \sum_{j=1}^J \lambda_{\min}(H_j) - \gamma' \geq -\frac{r_t}{J}(1 + h^{-1}(t)) - \sqrt{\frac{2t}{J}}\beta_t \quad (\text{S36})$$

$$\sum_{j=1}^J \lambda_{\min}(H_j) \geq -r_t(1 + h^{-1}(t)) - \sqrt{2tJ}\beta_t + J\gamma', \quad (\text{S37})$$

with probability at least $1 - \exp(-t)$. Here, $h(t) = \frac{1}{2}(t - \log(1 + t))$. The two additional model-specific parameters are β_t and r_t , which do not depend on J .

We additionally require, for $c \in (0, \sum_{j=1}^J \mathbb{E}\lambda_{\min}(H_j))$,

$$\sum_{j=1}^J \lambda_{\min}(H_j) \geq -r_t(1 + h^{-1}(t)) - \sqrt{2tJ}\beta_t + J\gamma' \geq c, \quad (\text{S38})$$

$$J\gamma' \geq c + r_t(1 + h^{-1}(t)) + \sqrt{2tJ}\beta_t. \quad (\text{S39})$$

It is therefore sufficient to have

$$J > \max \left\{ \frac{2r_t(1 + h^{-1}(t)) + 2c}{\gamma'}, \frac{2(2t\beta_t^2 + c)^2}{\gamma'^2} \right\} \geq \frac{r_t(1 + h^{-1}(t))}{\gamma'} + \sqrt{2tJ}\beta_t/\gamma' + c/\gamma' \quad (\text{S40})$$

for $\hat{H}(\theta)$ to be invertible and positive definite with minimum eigenvalue greater than or equal to c with probability at least $1 - \exp(-t)$. □

We now proceed with the proof of Theorem 2. In this analysis, we largely follow the proof of Theorem 6 in Roosta-Khorasani and Mahoney (2016). Define $\theta_\eta = \theta_m + \eta p_m$, where $p_m = -(H(\theta_m))^{-1}g(\theta_m)$. As in Roosta-Khorasani and Mahoney (2016), we want to show there is some iteration-independent $\tilde{\eta} > 0$ such that the Armijo condition

$$f(\theta_m + \eta p_m) \leq f(\theta_m) + \eta \beta p_m^T g(\theta_m), \quad (\text{S41})$$

holds for any $0 < \eta < \tilde{\eta}$ and some $\beta \in (0, 1)$. By an argument found in the beginning of the proof of Theorem 6 of Roosta-Khorasani and Mahoney (2016), we have that choosing J such that $\|\nabla_\theta \hat{\ell}(\theta_m) - \nabla_\theta \ell(\theta_m)\| \leq \epsilon$ and $\lambda_{\min}(H(\theta_m)) \geq c > 0$ for each m , yields

$$f(\theta_\eta) - f(\theta_m) \leq \eta p_m^T g(\theta_m) + \epsilon \eta \|p_m\| + \eta^2 \Gamma \|p_m\|^2 / 2, \quad (\text{S42})$$

with probability $1 - \delta/2$. From now on, we assume that we are on the success event of this high-probability statement. Consequently, we have

$$p_m^T g(\theta_m) = -p_m^T H(\theta_m) p_m \geq -c \|p_m\|^2, \quad (\text{S43})$$

and we can obtain a decrease in the objective. Substituting this into the previous expression,

$$f(\theta_\eta) - f(\theta_m) \leq -\eta p_m^T H(\theta_m) p_m + \epsilon \eta \|p_m\| + \eta^2 \Gamma \|p_m\|^2 / 2, \quad (\text{S44})$$

the Armijo condition becomes

$$-\eta p_m^T H(\theta_m) p_m + \epsilon \eta \|p_m\| + \eta^2 \Gamma \|p_m\|^2 / 2 \leq \eta \beta p_m^T g(\theta_m) = -\eta \beta p_m^T H(\theta_m) p_m \quad (\text{S45})$$

$$\epsilon \|p_m\| + \eta \Gamma \|p_m\|^2 / 2 \leq (1 - \beta) p_m^T H(\theta_m) p_m \quad (\text{S46})$$

$$\epsilon + \eta \Gamma \|p_m\| / 2 \leq c(1 - \beta) \|p_m\|. \quad (\text{S47})$$

This holds and guarantees an iteration-independent lower bound if

$$\eta \leq \frac{c(1 - \beta)}{\Gamma}, \quad \epsilon \leq \frac{c(1 - \beta)}{2\Gamma} \|g(\theta_m)\| \leq \frac{c(1 - \beta)}{2} \|p_m\|, \quad (\text{S48})$$

which is given by our choice of η . Now, first note that

$$\|g(\theta_m)\| - \|\nabla_\theta f(\theta_m)\| \leq \|g(\theta_m) - \nabla_\theta f(\theta_m)\| \leq \epsilon \implies \|\nabla_\theta f(\theta_m)\| \geq \|g(\theta_m)\| - \epsilon \quad (\text{S49})$$

and

$$\|\nabla_\theta f(\theta_m)\| - \|g(\theta_m)\| \leq \|\nabla_\theta f(\theta_m) - g(\theta_m)\| \leq \epsilon \implies \|g(\theta_m)\| \geq \|\nabla_\theta f(\theta_m)\| - \epsilon. \quad (\text{S50})$$

There are now two cases. If the algorithm terminates and $\|g(\theta_m)\| \leq \sigma\epsilon$, we can derive

$$\|\nabla_\theta f(\theta_m)\| \leq \|g(\theta_m)\| + \epsilon = \sigma\epsilon + \epsilon = (\sigma + 1)\epsilon. \quad (\text{S51})$$

If the algorithm does not terminate, then $\|g(\theta_m)\| > \sigma\epsilon$. Notice that

$$\epsilon \geq \|g(\theta_m) - \nabla_\theta f(\theta_m)\| \geq \|g(\theta_m)\| - \|\nabla_\theta f(\theta_m)\| \quad (\text{S52})$$

$$\|\nabla_\theta f(\theta_m)\| + \epsilon \geq \|g(\theta_m)\| \geq \sigma\epsilon \quad (\text{S53})$$

$$\|\nabla_\theta f(\theta_m)\| \geq \sigma\epsilon - \epsilon = (\sigma - 1)\epsilon \quad (\text{S54})$$

$$\frac{\|\nabla_{\theta} f(\theta_m)\|}{\sigma - 1} \geq \epsilon, \quad (\text{S55})$$

and now

$$\|\nabla_{\theta} f(\theta_m)\| - \epsilon \geq \|\nabla_{\theta} f(\theta_m)\| - \frac{\|\nabla_{\theta} f(\theta_m)\|}{\sigma - 1} \quad (\text{S56})$$

$$= \left(1 - \frac{1}{\sigma - 1}\right) \|\nabla_{\theta} f(\theta_m)\| \quad (\text{S57})$$

$$= \frac{\sigma - 2}{\sigma - 1} \|\nabla_{\theta} f(\theta_m)\| \quad (\text{S58})$$

$$\geq \frac{2}{3} \|\nabla_{\theta} f(\theta_m)\|. \quad (\text{S59})$$

Since $\|A^{-1}\| = 1/\sigma_{\min}(A)$,

$$p_m^T H(\theta_m) p_m = (- (H(\theta_m))^{-1} g(\theta_m))^T H(\theta_m) (- (H(\theta_m))^{-1} g(\theta_m)) \quad (\text{S60})$$

$$= g(\theta_m)^T (H(\theta_m))^{-1} g(\theta_m) \quad (\text{S61})$$

$$\geq \frac{1}{c} \|g(\theta_m)\|^2 \quad (\text{S62})$$

$$\geq \frac{1}{c} (\|\nabla_{\theta} f(\theta_m)\| - \epsilon)^2 \quad (\text{S63})$$

$$\geq \frac{4}{9c} \|\nabla_{\theta} f(\theta_m)\|^2. \quad (\text{S64})$$

From the assumption that f is γ -strongly convex, $\gamma I \preceq \nabla_{\theta}^2 - \ell \preceq \Gamma I$, by an implication of γ -strong convexity we have

$$f(\theta_m) - f(\theta^*) \leq \frac{1}{2\gamma} \|\nabla_{\theta} f(\theta_m)\|^2, \quad (\text{S65})$$

and we put together:

$$f(\theta_m) - f(\theta^*) \leq \frac{1}{2\gamma} \|\nabla_{\theta} f(\theta_m)\|^2 \leq \frac{9c}{4} \frac{1}{2\gamma} p_m^T H(\theta_m) p_m, \quad (\text{S66})$$

$$\frac{8\gamma}{9c} (f(\theta_m) - f(\theta^*)) \leq \frac{4}{9c} \|\nabla_{\theta} f(\theta_m)\|^2 \leq p_m^T H(\theta_m) p_m, \quad (\text{S67})$$

$$f(\theta_m) - f(\theta^*) \leq \frac{9c}{8\gamma} p_m^T H(\theta_m) p_m, \quad (\text{S68})$$

$$-\frac{8\gamma}{9c} (f(\theta_m) - f(\theta^*)) \geq -\frac{4}{9c} \|\nabla_{\theta} f(\theta_m)\|^2 \geq -p_m^T H(\theta_m) p_m. \quad (\text{S69})$$

From earlier, as the Armijo condition is fulfilled with our choice of η and ϵ ,

$$f(\theta_{m+1}) - f(\theta_m) \leq -\eta p_m^T H(\theta_m) p_m + \epsilon \eta \|p_m\| + \eta^2 \Gamma \|p_m\|^2 / 2 \quad (\text{S70})$$

$$\leq -\eta \beta p_m^T H(\theta_m) p_m \quad (\text{S71})$$

$$\leq -\eta \beta \frac{8\gamma}{9c} (f(\theta_m) - f(\theta^*)). \quad (\text{S72})$$

Therefore,

$$f(\theta_{m+1}) - f(\theta^*) = f(\theta_{m+1}) - f(\theta_m) + f(\theta_m) - f(\theta^*) \quad (\text{S73})$$

$$\leq f(\theta_m) - f(\theta^*) - \eta\beta \frac{8\gamma}{9c} (f(\theta_m) - f(\theta^*)) \quad (\text{S74})$$

$$= \left(1 - \eta\beta \frac{8\gamma}{9c}\right) (f(\theta_m) - f(\theta^*)). \quad (\text{S75})$$

S3 Feynman-Kac Models and Monte Carlo Approximations

In this section, we introduce the Feynman-Kac convention of Del Moral (2004) that has since become commonplace (Karjalainen et al., 2023) for the analysis of the particle filter. The mathematical formalization and notation introduced here will be adopted in the remainder of the analysis, in order to prove Theorems 3, 4, and 5. Let $(\eta_n)_{n=1}^N, (\pi_n)_{n=1}^N, (\rho_n)_{n=1}^N$ be sequences of probability measures on the state space $\mathcal{X}$. This is the sequence of prediction distributions $f_{X_n|Y_{1:n-1}}$, filtering distributions $f_{X_n|Y_{1:n}}$, and posterior distributions $f_{X_{1:n}|Y_{1:n}}$ that we seek to approximate with the particle filter. For any measurable bounded functional h , we adopt the following functional-analytic notation, borrowed from Del Moral (2004); Chopin and Papaspiliopoulos (2020); Karjalainen et al. (2023). We choose our specific choice of notation and definitions to be in line with that of Karjalainen et al. (2023).

Markov kernels and the process model: A Markov kernel M with source $\mathcal{X}_1$ and target $\mathcal{X}_2$ is a map $M : \mathcal{X}_1 \times \mathcal{B}(\mathcal{X}_2) \rightarrow [0, 1]$ such that for every set $A \in \mathcal{B}(\mathcal{X}_2)$ and every point $x \in \mathcal{X}_1$, the map $x \mapsto M(x, A)$ is a measurable function of x , and the map $A \mapsto M(x, A)$ is a probability measure on $\mathcal{X}_2$. The quantity $M(x, A)$ can be thought of as the probability of transitioning to the set A given that we are at the point x . If this yields a density, this then corresponds to the process density $f_{X_2|X_1}$ conditional on x and integrated over A .

Markov kernels and measures: For any measure η , any Markov kernel M on $\mathcal{X}$, any point $x \in \mathcal{X}$ and any measurable subset $A \subseteq \mathcal{X}$, let

$$\eta(h) = \int h d\eta = \int h(x)\eta(dx), \quad (\text{S76})$$

$$(\eta M)(A) = \int \eta(dx)M(x, A), \quad (\text{S77})$$

$$(Mh)(x) = \int M(x, dy)h(y). \quad (\text{S78})$$

Compositions of Markov kernels: The composition of a Markov kernel M_1 with another Markov kernel M_2 is another Markov kernel, given by

$$(M_1 M_2)(x, A) = \int M_1(x, dy)M_2(y, A). \quad (\text{S79})$$

Total variation distance: The total variation distance between two measures μ and ν on $\mathcal{X}$ is

$$\|\mu - \nu\|_{\text{TV}} = \sup_{\|h\|_{\infty} \leq 1/2} |\mu(h) - \nu(h)| = \sup_{\text{osc}(h) \leq 1} |\mu(h) - \nu(h)|. \quad (\text{S80})$$

Dobrushin contraction: The Dobrushin contraction coefficient β_{TV} of a Markov kernel M is given by

$$\beta_{\text{TV}}(M) = \sup_{x,y \in \mathcal{X}} \|M(x, \cdot) - M(y, \cdot)\|_{\text{TV}} = \sup_{\mu, \nu \in \mathcal{P}, \mu \neq \nu} \frac{\|\mu M - \nu M\|_{\text{TV}}}{\|\mu - \nu\|_{\text{TV}}}. \quad (\text{S81})$$

Potential functions and the measurement model: A potential function $G : \mathcal{X} \rightarrow [0, \infty)$ is a non-negative function of an element of the state space $x \in \mathcal{X}$. In our case, this corresponds to the measurement model, and in our previous notation is written as $g_{n,j} = f_{Y_n|X_n}(y_n^*|x_{n,j}^F; \theta) = G_n(x_{n,j}^F)$, where we suppress the dependence on θ for notational simplicity. Note that $G_n(\cdot) = f_{Y_n|X_n}(y_n^*|\cdot)$ is the conditional density of the observed measurement at time n , where we condition on the filtering particle $x_{n,j}^F$ as an element of the state space.

Feynman-Kac models: A Feynman-Kac model on $\mathcal{X}$ is a tuple $(\pi_0, (M_n)_{n=1}^N, (G_n)_{n=1}^N)$ of an initial probability measure on the state space π_0 , a sequence of transition kernels $(M_n)_{n=1}^N$, and a sequence of potential functions $(G_n)_{n=1}^N$. In the notation used in the main text, this corresponds to the starting distribution f_{X_0} , the sequence of transition densities $f_{X_n|X_{n-1}}$, and the measurement densities $f_{Y_n|X_n}$. This induces a set of mappings from the set of probability measures on $\mathcal{X}$ to itself, $\mathcal{P}(\mathcal{X}) \rightarrow \mathcal{P}(\mathcal{X})$, as follows:

- The update from the prediction to the filtering distributions is given by

$$\pi_n(dx) = \Psi_n(\eta_n)(dx) = \frac{G_n(x) \cdot \eta_n(dx)}{\eta_n(G_n)}. \quad (\text{S82})$$

- The map from the prediction distribution at timestep n to timestep $n+1$ is given by

$$\Phi_{n+1}(\eta_n) = \Psi_n(\eta_n)M_{n+1}. \quad (\text{S83})$$

- The composition of maps between prediction distributions yields the map from the prediction distribution at time k to the prediction distribution at time n where $k \leq n$,

$$\Phi_{k,n} = \Phi_n \circ \dots \circ \Phi_{k+1}. \quad (\text{S84})$$

The particle filter: The particle filter then yields a Monte Carlo approximation to the above Feynman-Kac model, via a sequence of mixture Dirac measures. When one resamples at every timestep, the prediction measure at timestep n is then given by

$$\eta_n^J = \frac{1}{J} \sum_{j=1}^J \delta_{x_{n,j}^P}, \quad (\text{S85})$$

and the filtering measure at timestep n is given by

$$\pi_n^J = \frac{\sum_{j=1}^J g_{n,j} \delta_{x_{n,j}^P}}{\sum_{j=1}^J g_{n,j}} \approx \frac{1}{J} \sum_{j=1}^J \delta_{x_{n,j}^F}. \quad (\text{S86})$$

In a slight abuse of notation, we will identify $x_{n,1:J}^P \equiv \eta_n^J$, and $x_{n,1:J}^F \equiv \pi_n^J$. As in Karjalainen et al. (2023), one can view this as an inhomogenous Markov process evolving on $\mathcal{X}^J$. The corresponding Markov transition kernel is then

$$\mathbf{M}_n(x_{n-1,1:J}^P, \cdot) = (\Phi_n(\eta_{n-1}^J))^{\otimes J} = \left(\Phi_n \left(\frac{1}{J} \sum_{j=1}^J \delta_{x_{n-1,j}^P} \right) \right)^{\otimes J}, \quad (\text{S87})$$

and the composition of Markov kernels on particles from timestep n to timestep $n+k$ is written

$$\mathbf{M}_{n,n+k} = \mathbf{M}_{n+1} \cdots \mathbf{M}_{n+k}. \quad (\text{S88})$$

One may wonder why Karjalainen et al. (2023) require this process to evolve on $\mathcal{X}^J$. This is because at every timestep n , we in fact draw $X_{n,j}^P | \{X_{n-1,1:J}^P = x_{n-1,1:J}^P\} \sim \mathbf{M}_n(x_{n-1,1:J}^P, \cdot) = \eta_0^{\otimes J} \mathbf{M}_{0,n}$ for $j = 1, \dots, J$.

Forgetting of the particle filter: The above formalization yields a result from Karjalainen et al. (2023) on the forgetting of the particle filter that we require for our analysis of the bias, variance, and error of MOP- α . That is, Karjalainen et al. (2023) show that

$$\beta_{\text{TV}}(\mathbf{M}_{n,n+k}) \leq (1 - \epsilon)^{\lfloor k/(O(\log J)) \rfloor}, \quad (\text{S89})$$

for some ϵ dependent on $\bar{G}, \underline{G}, \bar{M}, \underline{M}$ in Assumptions (A3) and (A2). As a result, the mixing time of the particle filter is only on the order of $O(\log(J))$ timesteps.

Equipped with the above formalisms and results, we are now in a position to provide guarantees on the performance of MOP- α itself.

S4 Proof of Theorem 3 via a Strong Law of Large Numbers for Triangular Arrays of Particles With Off-Parameter Resampling

Here, we prove Theorem 3 by deriving more general result, from which it will be clear that MOP- α targets the filtering distribution and is strongly consistent for the likelihood. We will prove a strong law of large numbers for triangular arrays of particles with off-parameter resampling, meaning that we resample the particles according to an arbitrary resampling rule that is not necessarily in proportion to the target distribution of interest. Using weights that encode the cumulative discrepancy between the resampling distribution and the target distribution (instead of resampling to equal weights, as in the basic particle filter) provides a sufficient correction to ensure almost sure convergence. We now introduce the precise definition of an off-parameter resampled particle filter.

Definition S1 (Off-Parameter Resampled Particle Filters). *An off-parameter resampled particle filter is a Monte Carlo approximation to a Feynman-Kac model $(\pi_0, (M_n)_{n=1}^N, (G_n)_{n=1}^N)$, where we inductively define given some $x_{n-1,j}^F, w_{n-1,j}^F$ comprising the filtering measure approximation π_n^J :*

1. *The prediction particles at timestep n , $x_{n,j}^P$, are drawn from $X_{n,j}^P \sim \pi_n^J M_n$, and the prediction weights are given by $w_{n,j}^P = w_{n-1,j}^F$.*

2. The prediction measure at timestep n is given by

$$\eta_n^J = \frac{\sum_{j=1}^J w_{n,j}^P \delta_{x_{n,j}^P}}{\sum_{j=1}^J w_{n,j}^P}. \quad (\text{S90})$$

3. The particles are resampled at every timestep n , yielding indices k_j according to some arbitrary probabilities $(p_{n,j})_{j=1}^J$.

4. The filtering measure at timestep n is given by

$$\pi_n^J = \frac{\sum_{j=1}^J \delta_{x_{n,j}^P} w_{n,j}^P g_{n,j}}{\sum_{j=1}^J w_{n,j}^P g_{n,j}} \approx \frac{\sum_{j=1}^J \delta_{x_{n,k_j}^P} w_{n,k_j}^P g_{n,k_j} / p_{n,k_j}}{\sum_{j=1}^J w_{n,k_j}^P g_{n,k_j} / p_{n,k_j}} = \frac{\sum_{j=1}^J w_{n,j}^F \delta_{x_{n,j}^F}}{\sum_{j=1}^J w_{n,j}^F}. \quad (\text{S91})$$

5. The prediction weights at the next timestep are given by $w_{n+1,j}^P = w_{n,j}^F = w_{n,k_j}^P g_{n,k_j} / p_{n,k_j}$.

To prove that the weight correction is sufficient for almost sure convergence, we first introduce what it means for a triangular array of particles to *target* a given target distribution.

Definition S2 (Targeting). A pair of random vectors (X, W) drawn from some measure g **targets** another measure π if, for any measurable and bounded functional h ,

$$E_g[h(X) \cdot W] = E_\pi[h(X)]. \quad (\text{S92})$$

A particle approximation is a triangular array of pairs of random vectors $(X_j^J, W_j^J), j = 1, 2, \dots, J$. The particle approximation **targets** π if, for any measurable and bounded functional h ,

$$\frac{\sum_{j=1}^J h(X_j^J) W_j^J}{\sum_{j=1}^J W_j^J} \xrightarrow{\text{a.s.}} E_\pi[h(X)] \quad (\text{S93})$$

as $J \rightarrow \infty$.

Chopin (2004) asserted without proof that common particle filter algorithms target the filtering distribution in this sense, while Chopin and Papaspiliopoulos (2020) proved a related result assuming bounded densities. We follow a similar approach to Chopin and Papaspiliopoulos (2020), based on showing strong laws of large numbers for triangular arrays, noting that triangular array strong laws do not generally hold without an additional regularity condition such as boundedness.

In order to prove the consistency of our variation on the particle filter, we now present three helper lemmas. The first follows from standard importance sampling arguments, the second from integrating out the marginal, and the third from Bayes' theorem. We state Lemma S5 assuming multinomial resampling, which is convenient for the proof though other resampling strategies may be preferable in practice.

Lemma S5 (Change of Weight Measure). Suppose that $\{(\tilde{X}_j^J, U_j^J), j = 1, \dots, J\}$ targets f_X . Now, let $\{(Y_j^J, V_j^J), j = 1, \dots, J\}$ be a multinomial sample with indices k_j drawn from $\{(\tilde{X}_j^J, U_j^J)\}$ where $(\tilde{X}_j^J, U_j^J)$ is represented, on average, proportional to $\pi_j^J J$ times. Write

$$(Y_j^J, V_j^J) = (\tilde{X}_{k_j}^J, U_{k_j}^J / \pi_{k_j}^J). \quad (\text{S94})$$

If the importance sampling weights U_j / π_j are bounded, then $\{(Y_j^J, V_j^J), j = 1, \dots, J\}$ targets f_X .

Proof. Note that as the Y_j^J are a subsample from X_j^J , h can be a function of Y as well as it is one for X . We then expand

$$\frac{\sum_j h(Y_j^J) V_j^J}{\sum_j V_j^J} = \frac{\sum_j h(\tilde{X}_{k_j}^J) \frac{U_{k_j}^J}{\pi_{k_j}^J}}{\sum_j \frac{U_{k_j}^J}{\pi_{k_j}^J}}. \quad (\text{S95})$$

By hypothesis,

$$\frac{\sum_j h(\tilde{X}_j^J) U_j^J}{\sum_j U_j^J} \xrightarrow{a.s.} \mathbb{E}_{f_X} [f(X)]. \quad (\text{S96})$$

We want to show

$$\frac{\sum_j h(X_{k_j}^J) \frac{U_{k_j}^J}{\pi_{k_j}^J}}{\sum_j \frac{U_{k_j}^J}{\pi_{k_j}^J}} - \frac{\sum_j h(\tilde{X}_j^J) U_j^J}{\sum_j U_j^J} \xrightarrow{a.s.} 0. \quad (\text{S97})$$

For this, it is sufficient to show that

$$\sum_j h(\tilde{X}_{k_j}^J) \frac{U_{k_j}^J}{\pi_{k_j}^J} - \sum_j h(\tilde{X}_j^J) \frac{U_j^J}{\pi_j^J} \xrightarrow{a.s.} 0 \quad (\text{S98})$$

since an application of this result with $h(x) = 1$ provides almost sure convergence of the denominator in (S97). Write $g(\tilde{X}_j^J) = h(\tilde{X}_j^J) \frac{U_{k_j}^J}{\pi_{k_j}^J}$. We therefore need to show that

$$\sum_j Z_j^J := \sum_j \left(g(\tilde{X}_{k_j}^J) - g(\tilde{X}_j^J) \pi_j^J \right) \xrightarrow{a.s.} 0. \quad (\text{S99})$$

Because the functional h and importance sampling weights $u_{k_j}^J/\pi_{k_j}^J$ are bounded, we have that $\mathbb{E}[(Z_j^J)^4] < \infty$. We can then follow the argument of Chopin and Papaspiliopoulos (2020) from this point on, where noting that the Z_j^J are conditionally independent given the $\tilde{X}_j^J$ and U_j^J ,

$$\begin{aligned} \mathbb{E} \left[\left(\sum_j Z_j^J \right)^4 \middle| (\tilde{X}_j^J, U_j^J)_{j=1}^J \right] &= J \mathbb{E} \left[(Z_1^J)^4 | (\tilde{X}_1^J, U_1^J) \right] \\ &\quad + 3J(J-1) \left(\mathbb{E}[(Z_j^J)^2 | (\tilde{X}_j^J, U_j^J)_{j=1}^J] \right) \end{aligned} \quad (\text{S100})$$

$$\leq C J^2, \quad (\text{S101})$$

for some $C > 0$. Taking expectations on both sides yields

$$\mathbb{E} \left[\left(\sum_j Z_j^J \right)^4 \right] \leq C J^2 \quad (\text{S102})$$

by the tower property. Now by Markov's inequality,

$$\mathbb{P} \left(\left| \frac{1}{J} \sum_{j=1}^J Z_j^J \right| > \epsilon \right) \leq \frac{1}{\epsilon^4 J^4} \mathbb{E} \left[\left(\sum_{j=1}^J Z_j^J \right)^4 \right] \leq \frac{C}{\epsilon^4 J^2}, \quad (\text{S103})$$

and as these terms are summable we can apply Borel-Cantelli to conclude that these deviations happen only finitely often for every $\epsilon > 0$, giving us the almost-sure convergence for

$$\sum_j h(X_{k_j}^J) \frac{u_{k_j}^J}{\pi_{k_j}^J} - \sum_j h(X_j^J) \frac{u_j^J}{\pi_j^J} \xrightarrow{a.s.} 0. \quad (\text{S104})$$

Similarly, we also have that

$$\sum_j \frac{u_j^J}{\pi_j^J} \pi_j^J - \sum_j \frac{u_{k_j}^J}{\pi_{k_j}^J} \pi_{k_j}^J \xrightarrow{a.s.} 0, \quad (\text{S105})$$

and the result is proved. $\square$

Remark: Note that Lemma S5 permits $\pi_{1:J}$ to depend on $\{(X_j, u_j)\}$ as long as the resampling is carried out independently of $\{(X_j, u_j)\}$, conditional on $\pi_{1:J}$.

Lemma S6 (Particle Marginals). *Suppose that $\{(\tilde{X}_j^J, U_j^J), j = 1, \dots, J\}$ targets f_X . Also suppose that $\tilde{Z}_j^J \sim f_{Z|X}(\cdot | \tilde{X}_j^J)$ where $f_{Z|X}$ is a conditional probability density function corresponding to a joint density $f_{X,Z}$ with marginal densities f_X and f_Z . Then, if the U_j^J are bounded, $\{(\tilde{Z}_j^J, U_j^J)\}$ targets f_Z .*

Proof. We want to show that, for any measurable bounded h ,

$$\frac{\sum_j h(\tilde{Z}_j^J) U_j^J}{\sum_j U_j^J} \xrightarrow{a.s.} \mathbb{E}_{f_Z}[h(Z)] = \mathbb{E}_{f_X}[\mathbb{E}_{f_{Z|X}}[h(Z)|X]]. \quad (\text{S106})$$

By assumption, for any measurable and bounded functional g with domain $\mathcal{X}$,

$$\frac{\sum_j g(\tilde{X}_j^J) U_j^J}{\sum_j U_j^J} \xrightarrow{a.s.} \mathbb{E}_{f_X}[g(X)]. \quad (\text{S107})$$

Let $\bar{U}_j^J = \frac{J U_j^J}{\sum_j U_j^J}$. Examine the numerator and denominator of the quantity

$$\frac{J^{-1} \sum_j h(\tilde{Z}_j^J) U_j^J}{J^{-1} \sum_j U_j^J} = \frac{J^{-1} \sum_j h(\tilde{Z}_j^J) \bar{U}_j^J}{J^{-1} \sum_j \bar{U}_j^J}. \quad (\text{S108})$$

The denominator converges to 1 almost surely. The numerator, on the other hand, is

$$\frac{1}{J} \sum_j h(\tilde{Z}_j^J) \bar{U}_j^J, \quad (\text{S109})$$

and by the same fourth moment argument to the above lemma, it converges almost surely to the limit of its expectation,

$$\lim_{J \rightarrow \infty} \mathbb{E} \left[\frac{1}{J} \sum_j h(\tilde{Z}_j^J) \bar{U}_j^J \right] = \lim_{J \rightarrow \infty} \mathbb{E} \left[\frac{1}{J} \sum_j \mathbb{E} [h(\tilde{Z}_j^J) \bar{U}_j^J | \tilde{X}_j^J, \bar{U}_j^J] \right] \quad (\text{S110})$$

$$= \lim_{J \rightarrow \infty} \mathbb{E} \left[\frac{1}{J} \sum_j \mathbb{E} [h(Z) | X = \tilde{X}_j^J] \bar{U}_j^J \right]. \quad (\text{S111})$$

Applying (S107) with $g(x) = \mathbb{E}[h(Z)|X = x]$, the average on the right hand side converges almost surely to $\mathbb{E}\{\mathbb{E}[h(Z)|X]\} = \mathbb{E}[h(Z)]$. It remains to swap the limit and expectations. We can do so with the bounded convergence theorem, and therefore obtain

$$\frac{1}{J} \sum_j h(\tilde{Z}_j^J) \bar{U}_j^J \xrightarrow{a.s.} \mathbb{E}_{f_Z}[h(Z)]. \quad (\text{S112})$$

□

Lemma S7 (Particle Posteriors). *Suppose that $\{(X_j^J, U_j^J), j = 1, \dots, J\}$ targets f_X . Also suppose that $(X_j'^J, U_j'^J) = (X_j^J, U_j^J f_{Z|X}(z^*|X_j^J))$. Then, if $U_j^J f_{Z|X}(z^*|X_j^J)$ and $U_j^J f_{Z|X}(z^*|X_j^J)/f_Z(z^*)$ are bounded, $\{(X_j'^J, U_j'^J)\}$ targets $f_{X|Z}(\cdot|z^*)$.*

Proof. Again, we want to show that

$$\frac{\sum_j h(X_j^J) \cdot U_j^J \cdot f_{Z|X}(z^*|X_j^J)}{\sum_j U_j^J \cdot f_{Z|X}(z^*|X_j^J)} \xrightarrow{a.s.} \mathbb{E}_{f_{X|Z}}[h(X)|z^*]. \quad (\text{S113})$$

We already have that for any measurable bounded g ,

$$\frac{\sum_j g(X_j^J) U_j^J}{\sum_j U_j^J} \xrightarrow{a.s.} \mathbb{E}_{f_X}[g(X)]. \quad (\text{S114})$$

Consider the following:

$$\frac{J^{-1} \sum_j h(X_j^J) f_{Z|X}(z^*|X_j^J) U_j^J}{J^{-1} \sum_j U_j^J} \times \left(\frac{J^{-1} \sum_j f_{Z|X}(z^*|X_j^J) U_j^J}{J^{-1} \sum_j U_j^J} \right)^{-1}. \quad (\text{S115})$$

We will apply Equation (S114) to the numerator and the denominator in the ratio above individually. The numerator converges to

$$\frac{J^{-1} \sum_j h(X_j^J) f_{Z|X}(z^*|X_j^J) U_j^J}{J^{-1} \sum_j U_j^J} \xrightarrow{a.s.} \mathbb{E}_{f_X}[h(X) f_{Z|X}(z^*|X)], \quad (\text{S116})$$

while the reciprocal of the denominator converges to

$$\frac{J^{-1} \sum_j f_{Z|X}(z^*|X_j^J) U_j^J}{J^{-1} \sum_j U_j^J} \xrightarrow{a.s.} \mathbb{E}_{f_X}[f_{Z|X}(z^*|X)] = f_Z(z^*). \quad (\text{S117})$$

Now we take advantage of the identities

$$\frac{\mathbb{E}_{f_X}[h(X) f_{Z|X}(z^*|X)]}{f_Z(z^*)} = \mathbb{E}_{f_X} \left[\frac{h(X) f_{Z|X}(z^*|X)}{f_Z(z^*)} \right] \quad (\text{S118})$$

$$= \mathbb{E}_{f_X} \left[h(X) \frac{f_{X|Z}(X|z^*)}{f_X(X)} \right] = \mathbb{E}_{f_{X|Z}}[h(X)|z^*], \quad (\text{S119})$$

to give the desired result.

□

Theorem S1 (Off-Parameter Resampled Particle Filters Target the Filtering Distribution). *The off-parameter resampled particle filter as outlined in Definition S1 targets the filtering distribution.*

Proof. Recursively applying Lemmas S5, S6, and S7, we obtain that the off-parameter resampled particle filter targets the posterior. Specifically, suppose inductively that $\{(X_{n-1,j}^F, w_{n-1,j}^F)\}$ targets π_{n-1} . Then, Lemma S6 tells us that $\{(X_{n,j}^P, w_{n,j}^P)\}$ targets η_n . Lemma S7 tells us that $\{(X_{n,j}^P, w_{n,j}^P g_{n,j}^\theta)\}$ therefore targets π_n . Lemma S5 guarantees that the resampling rule, given by

$$(X_{n,j}^F, w_{n,j}^F) = (X_{n,k_j}^P, w_{n,k_j}^P g_{n,k_j} / p_{n,k_j}), \quad (\text{S120})$$

with resampling probabilities proportional to $p_{n,j}$, therefore also targets π_n . $\square$

Proposition S1 (MOP-1 Targets the Filtering Distribution). *When $\alpha = 1$ or $\phi = \theta$, MOP- α targets the filtering distribution.*

Proof. When $\theta = \phi$, regardless of the value of α , the ratio $\frac{g_{n,j}^\theta}{g_{n,j}^\phi} = 1$, and this reduces to the vanilla particle filter estimate.

When $\alpha = 1$, and $\theta \neq \phi$, the proof is identical to that of S1. Recursively applying Lemmas S5, S6, and S7, we obtain that the MOP-1 filter targets the posterior. Specifically, suppose inductively that $\{(X_{n-1,j}^{F,\theta}, w_{n-1,j}^{F,\theta})\}$ is properly weighted for $f_{X_{n-1}|Y_{1:n-1}}(x_{n-1}|y_{1:n-1}^*; \theta)$. Then, Lemma S6 tells us that $\{(X_{n,j}^{P,\theta}, w_{n,j}^{P,\theta})\}$ targets $f_{X_n|Y_{1:n-1}}(x_n|y_{1:n-1}^*; \theta)$. Lemma S7 tells us that $\{(X_{n,j}^{P,\theta}, w_{n,j}^{P,\theta} g_{n,j}^\theta)\}$ therefore targets $f_{X_n|Y_{1:n}}(x_n|y_{1:n}^*; \theta)$. Lemma S5 guarantees that the resampling rule, given by

$$(X_{n,j}^{F,\theta}, w_{n,j}^{F,\theta}) = (X_{n,k_j}^{P,\theta}, w_{n,k_j}^{P,\theta} g_{n,k_j}^\theta / g_{n,k_j}^\phi), \quad (\text{S121})$$

with resampling weights proportional to $g_{n,j}^\phi$, therefore also targets $f_{X_n|Y_{1:n}}(x_n|y_{1:n}^*; \theta)$. $\square$

This has addressed filtering, but not quite yet the likelihood evaluation. For this we use the following lemma.

Lemma S8 (Likelihood Proper Weighting). *$f_{Y_n|Y_{1:n-1}}(y_n^*|y_{1:n-1}^*; \theta)$ is strongly consistently estimated by either the before-resampling estimate,*

$$L_n^{B,\theta} = \frac{\sum_{j=1}^J g_{n,j}^\theta w_{n,j}^{P,\theta}}{\sum_{j=1}^J w_{n,j}^{P,\theta}}, \quad (\text{S122})$$

or by the after-resampling estimate,

$$L_n^{A,\theta} = L_n^\phi \frac{\sum_{j=1}^J w_{n,j}^{F,\theta}}{\sum_{j=1}^J w_{n,j}^{P,\theta}}. \quad (\text{S123})$$

where L_n^ϕ is as defined in the various algorithms.

Here, (S122) is a direct consequence of our earlier result that $\{(X_{n,j}^{P,\theta}, w_{n,j}^{P,\theta})\}$ targets $f_{X_n|Y_{1:n-1}}(x_n|y_{1:n-1}^*; \theta)$. To see (S123), we write the numerator of (S5) as

$$L_n^\phi \sum_{j=1}^J \left[\frac{g_{n,j}^\theta}{g_{n,j}^\phi} w_{n,j}^{P,\theta} \right] \frac{g_{n,j}^\phi}{L_n^\phi} = L_n^\phi \sum_{j=1}^J w_{n,j}^{FC,\theta} \frac{g_{n,j}^\phi}{L_n^\phi}. \quad (\text{S124})$$

Using Lemma S5, we resample according to probabilities $\frac{g_{n,j}^\phi}{L_n^\phi}$ to see this is properly estimated by

$$L_n^\phi \sum_{j=1}^J w_{n,j}^{F,\theta}, \quad (\text{S125})$$

from which we obtain (S123). Using Lemma S8, we obtain a likelihood estimate,

$$L^{A,\theta} = \prod_{n=1}^N \left(L_n^\phi \frac{\sum_{j=1}^J w_{n,j}^{F,\theta}}{\sum_{j=1}^J w_{n,j}^{P,\theta}} \right). \quad (\text{S126})$$

Since $w_{n,j}^{F,\theta} = w_{n+1,j}^{P,\theta}$, this is a telescoping product. The remaining terms are $\sum_{j=1}^J w_{0,j}^{P,\theta} = J$ on the denominator and $\sum_{j=1}^J w_{N,j}^{F,\theta}$ on the numerator. This derives the MOP-1 likelihood estimates.

$L^{B,\theta}$ should generally be preferred in practice, since there is no reason to include the extra variability from resampling when calculating the conditional log likelihood, but it lacks the nice telescoping product that lets us derive exact expressions for the gradient in Theorem 1.

S5 Proof of Theorem 4 via Consistency of Off-Parameter Resampled Gradient Estimates

We now provide a result showing that under sufficient regularity conditions, one can interchange the order of differentiation and expectation of functionals of off-parameter resampled particle estimates, showing that if an off-parameter resampled particle estimate is consistent for some estimand, its derivative is consistent for the derivative of the estimand as well.

One may wonder why this may be necessary, given that we have already shown strong consistency for measurable bounded functionals in Theorem S1. If we require the gradient to be bounded as well, then the gradient is also a bounded functional. The answer is that the strong consistency is for expectations under the filtering distribution, so Theorem S1 only establishes strong consistency of the gradient of the estimate to its expectation under the filtering distribution, $\nabla_\theta \eta_n^J(h_\theta) \xrightarrow{a.s.} \eta_n(\nabla_\theta h_\theta)$, which is not a-priori the gradient of the estimand $\nabla_\theta \eta_n(h_\theta)$. We require an interchange of the derivative and expectation, which we show is possible when the particle estimates have two continuous uniformly bounded derivatives over all J below.

Theorem S2 (Off-Parameter Particle Filters Yield Strongly Consistent Estimates of Derivatives of Functionals). *Let $h_\theta : \mathcal{X} \rightarrow \mathbb{R}$ be a measurable bounded functional of particles, where $\eta_n^J(h_\theta)$ has two continuous derivatives uniformly bounded over all J by H^* for almost every $\omega \in \Omega$ and $\theta \in \Theta$. If it holds that $\eta_n^J(h_\theta) \xrightarrow{a.s.} \eta(h_\theta) = h_\theta^*$, then we also have that $\nabla_\theta \eta_n^J(h_\theta) \xrightarrow{a.s.} \eta_n(\nabla_\theta h_\theta) = \nabla_\theta \eta_n(h_\theta) = \nabla_\theta h_\theta^*$.*

Proof. Fix $\omega \in \Omega$. The sequence $(\nabla_\theta \eta_n^J(h_\theta)(\omega))_{J \in \mathbb{N}}$ is uniformly bounded over all J , by assumption. The sequence is also uniformly equicontinuous. To see this, by assumption, the second derivative of $\eta_n^J(h_\theta)(\omega)|_{\theta=\theta'}$ is also uniformly bounded over all J for almost every $\omega \in \Omega$ and every $\theta' \in \Theta$. A set of functions with derivatives bounded by the same constant is uniformly Lipschitz, and therefore uniformly equicontinuous. So the sequence $(\eta_n^J(h_\theta)(\omega))_{J \in \mathbb{N}}$ is uniformly equicontinuous over θ for

almost every $\omega \in \Omega$. Explicitly, for almost every $\omega \in \Omega$ and every $\epsilon > 0$, there exists some $\delta(\omega) > 0$ such that for every $\|\theta - \theta'\|_\infty < \delta$ and every $J \in \mathbb{N}$ we have that

$$\|\eta_n^J(h_\theta)(\omega) - \eta_n^J(h_{\theta'})(\omega)\|_\infty < \epsilon. \quad (\text{S127})$$

Then, by Arzela-Ascoli, there exists a uniformly convergent subsequence. We claim that there is only one subsequential limit. When the gradient is bounded, we can treat the gradient as a bounded functional. So by Theorem S1 the sequence $(\eta_n^J(h_\theta)(\omega))_{J \in \mathbb{N}}$ converges pointwise for $\theta = \phi$ and almost every $\omega \in \Omega$, and there is therefore only one subsequential limit. The sequence therefore converges uniformly to its limit $\lim_{J \rightarrow \infty} \eta_n^J(h_\theta)(\omega)$. Therefore, with uniform convergence for the derivatives established, we can swap the limit and derivative, and obtain that for almost every $\omega \in \Omega$,

$$\lim_{J \rightarrow \infty} \eta_n^J(h_\theta)(\omega) = \nabla_\theta \lim_{J \rightarrow \infty} \eta_n^J(h_\theta)(\omega). \quad (\text{S128})$$

Again from Theorem S1, we know that

$$\eta_n^J(h_\theta)(\omega) \rightarrow \eta_n(h_\theta) = h_\theta^* \quad (\text{S129})$$

for almost every $\omega \in \Omega$. We then have that for almost every $\omega \in \Omega$,

$$\lim_{J \rightarrow \infty} \eta_n^J(h_\theta)(\omega) = \nabla_\theta \eta_n(h_\theta) = \nabla_\theta h_\theta^*, \quad (\text{S130})$$

as we wanted. $\square$

Theorem 4 is proved as a corollary of Theorem S2, where we apply $\eta_n^J(h_\theta) = \hat{\mathcal{L}}_J^1(\theta)$, $\eta_n(h_\theta) = \mathcal{L}(\theta) = h_\theta^*$, and then use the continuous mapping theorem.

S6 Proof of Theorem 5 (Bias-Variance Analysis)

In this section, we prove various rates on the bias, variance, and MSE of DMOP- α . The earlier truncation argument is retained below for comparison. A differentiable particle includes its location and tangent, so the particle-system forgetting result applies to the score functionals carried by that particle. With the manuscript's convention that $\lesssim$ suppresses the particle-system mixing-time factor, this gives the N/J -times-an- α -factor variance bound in Theorem S5. Section S7 gives the complete proof, together with a concise comparison with the path-space endpoint bounds.

First, we note the following relation between the Dobrushin contraction coefficient and the alpha-mixing coefficients in our context below.

Lemma S9. *Setting X to be the particle collection at time m , and Y at time n , we have that the alpha mixing coefficients,*

$$\int |f_{XY}(x, y) - f_X(x) f_Y(y)| dx dy < \alpha, \quad (\text{S131})$$

are bounded by the Dobrushin coefficient, i.e we have

$$\int |f_{Y|X}(y|x_1) - f_{Y|X}(y|x_2)| dy < \alpha \quad (\text{S132})$$

for all x_1, x_2 .

Proof. We rewrite the alpha-mixing assertion as

$$\int \int |f_{Y|X}(y|x) - f_Y(y)| dy f_X(x) dx. \quad (\text{S133})$$

We claim that the Dobrushin coefficient implies

$$\int |f_{Y|X}(y|x_1) - f_Y(y)| dy < \alpha \quad (\text{S134})$$

for all x_1 . This is shown as follows:

$$\int |f_{Y|X}(y|x) - f_Y(y)| dy = \int \left| \int [f_{Y|X}(y|x_1) - f_{Y|X}(y|x)] f_X(x) dx \right| dy \quad (\text{S135})$$

$$< \int \int |f_{Y|X}(y|x_1) - f_{Y|X}(y|x)| dy f_X(x) dx \quad (\text{S136})$$

$$< \int \alpha f_X(x) dx \quad (\text{S137})$$

$$= \alpha \quad (\text{S138})$$

We then have the desired result. $\square$

As a warm-up, we consider a variance bound for DMOP-0. Note that we made Assumptions (A2) and (A3) to leverage the results from Karjalainen et al. (2023) on the forgetting of the particle filter. This is used below to recover the $O(N/J)$ variance bound for DMOP-0 and a fixed-lag comparison. The final active result in Section S7 gives variance $O\{N/[J(1-\alpha)^2]\}$ for fixed $\alpha < 1$ and $O(N^3/J)$ at $\alpha = 1$, up to the common parameter and mixing-time factors. We start by decomposing the MOP- α estimator as follows, for the case $\theta = \phi$,

$$\hat{\mathcal{L}}(\theta) := \prod_{n=1}^N L_n^{A,\theta,\alpha} = \prod_{n=1}^N L_n^\phi \cdot \frac{\sum_{j=1}^J w_{n,j}^{F,\theta}}{\sum_{j=1}^J w_{n,j}^{P,\theta}} = \prod_{n=1}^N L_n^{A,\theta,\alpha} = \prod_{n=1}^N L_n^\phi \cdot \frac{\sum_{j=1}^J w_{n,j}^{F,\theta}}{\sum_{j=1}^J (w_{n-1,j}^{F,\theta})^\alpha}. \quad (\text{S139})$$

Now, to illustrate the proof strategy for the general case in an easier context, we first analyze the special case when $\alpha = 0$. We now prove the variance bound for this case, presented below. To our knowledge, this result for the variance of the gradient estimate of Naesseth et al. (2018) is novel.

Theorem S3 (Variance of MOP-0 Gradient Estimate). *The variance of DMOP-0, i.e. the algorithm of Naesseth et al. (2018), is*

$$\text{Var}(\nabla_\theta \hat{\ell}^0(\theta)) \lesssim \frac{Np G'(\theta)^2}{J}. \quad (\text{S140})$$

Proof. When $\alpha = 0$, we have

$$\hat{\mathcal{L}}^0(\theta) := \prod_{n=1}^N L_n^{A,\theta,0} = \prod_{n=1}^N L_n^\phi \cdot \frac{\sum_{j=1}^J w_{n,j}^{F,\theta,0}}{\sum_{j=1}^J w_{n,j}^{P,\theta,0}} \quad (\text{S141})$$

$$= \prod_{n=1}^N L_n^\phi \cdot \frac{1}{J} \sum_{j=1}^J s_{n,j} = \prod_{n=1}^N L_n^\phi \cdot \frac{1}{J} \sum_{j=1}^J \frac{f_{Y_n|X_n}(y_n^*|x_{n,j}^{P,\theta}; \theta)}{f_{Y_n|X_n}(y_n^*|x_{n,j}^{P,\phi}; \phi)}. \quad (\text{S142})$$

Similarly to the proof of Lemma S2, we define

$$s_{n,j} = \frac{f_{Y_n|X_n}(y_n^*|x_{n,j}^{P,\theta}; \theta)}{f_{Y_n|X_n}(y_n^*|x_{n,j}^{P,\phi}; \phi)}, \quad (\text{S143})$$

and from the exact same arguments conclude that its gradient when $\theta = \phi$ is therefore

$$\nabla_{\theta} \hat{\ell}^0(\theta) := \sum_{n=1}^N \nabla_{\theta} \log \left(L_n^{\phi} \frac{1}{J} \sum_{j=1}^J s_{n,j} \right) \quad (\text{S144})$$

$$= \sum_{n=1}^N \frac{\nabla_{\theta} \left(L_n^{\phi} \frac{1}{J} \sum_{j=1}^J s_{n,j} \right)}{\left(L_n^{\phi} \frac{1}{J} \sum_{j=1}^J s_{n,j} \right)} \quad (\text{S145})$$

$$= \sum_{n=1}^N \frac{\sum_{j=1}^J \nabla_{\theta} s_{n,j}}{\sum_{j=1}^J s_{n,j}} \quad (\text{S146})$$

$$= \sum_{n=1}^N \frac{1}{J} \sum_{j=1}^J \frac{\nabla_{\theta} f_{Y_n|X_n}(y_n^*|x_{n,j}^{F,\theta}; \theta)}{f_{Y_n|X_n}(y_n^*|x_{n,j}^{F,\phi}; \phi)} \quad (\text{S147})$$

$$= \frac{1}{J} \sum_{n=1}^N \sum_{j=1}^J \nabla_{\theta} \log \left(f_{Y_n|X_n}(y_n^*|x_{n,j}^{F,\theta}; \theta) \right), \quad (\text{S148})$$

where we use the derivative of the logarithm, observe that $\sum_{j=1}^J s_{n,j} = J$ when $\theta = \phi$, and use the derivative of the logarithm where $\theta = \phi$ again. We do this to establish that the gradient of the log-likelihood estimate is given by the sum of terms over all N and J :

$$\nabla_{\theta} \hat{\ell}^0(\theta) = \frac{1}{J} \sum_{n=1}^N \sum_{j=1}^J \nabla_{\theta} \log \left(f_{Y_n|X_n}(y_n^*|x_{n,j}^{F,\theta}; \theta) \right). \quad (\text{S149})$$

Therefore, for a given θ_i in the parameter vector θ ,

$$\text{Var} \left(\frac{\partial}{\partial \theta_i} \hat{\ell}^0(\theta) \right) = \frac{1}{J^2} \text{Var} \left(\sum_{n=1}^N \sum_{j=1}^J \frac{\partial}{\partial \theta_i} \log \left(f_{Y_n|X_n}(y_n^*|x_{n,j}^{F,\theta}; \theta) \right) \right) \quad (\text{S150})$$

$$\begin{aligned} &= \frac{1}{J^2} \sum_{n=1}^N \text{Var} \left(\sum_{j=1}^J \frac{\partial}{\partial \theta_i} \log \left(f_{Y_n|X_n}(y_n^*|x_{n,j}^{F,\theta}; \theta) \right) \right) \\ &\quad + 2 \sum_{m < n} \text{Cov} \left(\frac{1}{J} \sum_{j=1}^J \frac{\partial}{\partial \theta_i} \log \left(f_{Y_m|X_m}(y_m^*|x_{m,j}^{F,\theta}; \theta) \right), \frac{1}{J} \sum_{j=1}^J \frac{\partial}{\partial \theta_i} \log \left(f_{Y_n|X_n}(y_n^*|x_{n,j}^{F,\theta}; \theta) \right) \right) \\ &= \sum_{n=1}^N \text{Var} \left(\frac{\partial}{\partial \theta_i} \hat{\ell}_n^0(\theta) \right) + 2 \sum_{m < n} \text{Cov} \left(\frac{\partial}{\partial \theta_i} \hat{\ell}_m^0(\theta), \frac{\partial}{\partial \theta_i} \hat{\ell}_n^0(\theta) \right). \end{aligned} \quad (\text{S151})$$

Here, we use Assumptions (A2) and (A3) that ensure strong mixing. We know from Theorem 3.1 of Karjalainen et al. (2023) that when $\mathbf{M}_{n,n+k}$ is the k -step Markov operator from timestep n and

$\beta_{\text{TV}}(M) = \sup_{x,y \in E} \|M(x, \cdot) - M(y, \cdot)\|_{\text{TV}} = \sup_{\mu, \nu \in \mathcal{P}, \mu \neq \nu} \frac{\|\mu M - \nu M\|_{\text{TV}}}{\|\mu - \nu\|_{\text{TV}}}$ is the Dobrushin contraction coefficient of a Markov operator,

$$\beta_{\text{TV}}(\mathbf{M}_{n,n+k}) \leq (1 - \epsilon)^{\lfloor k/(c \log(J)) \rfloor}, \quad (\text{S152})$$

i.e. the mixing time of the particle filter is $O(\log(J))$, where ϵ and c depend on $\bar{M}, \underline{M}, \bar{G}, \underline{G}$ in (A2) and (A3).

The precise prediction-to-filter L^4 bridge and the passage from Theorem 3.1 of Karjalainen et al. (2023) to the score mixing coefficients are written in Equations (S287)– (S303) below. At $\alpha = 0$, they give

$$\left\| \frac{\partial}{\partial \theta_i} \hat{\ell}_n^0(\theta) - \mathbb{E} \left[\frac{\partial}{\partial \theta_i} \hat{\ell}_n^0(\theta) \right] \right\|_{L^4} \leq \frac{CG'(\theta)}{\sqrt{J}},$$

$$a(k) \leq (1 - \epsilon)^{\lfloor (k-1)/[c\{1+\log J\}] \rfloor}. \quad (\text{S153})$$

Davydov's inequality must be applied to the centered variables. Since covariance is unchanged by centering, Equation (S153) gives

$$\text{Cov} \left(\frac{\partial}{\partial \theta_i} \hat{\ell}_m^0(\theta), \frac{\partial}{\partial \theta_i} \hat{\ell}_n^0(\theta) \right) \leq (1 - \epsilon)^{\frac{1}{2} \lfloor \frac{|n-m|-1}{c\{1+\log(J)\}} \rfloor} \frac{G'(\theta)^2}{J}. \quad (\text{S154})$$

Putting it all together, we see that

$$\text{Var} \left(\frac{\partial}{\partial \theta_i} \hat{\ell}^\alpha(\theta) \right) = \sum_{n=1}^N \text{Var} \left(\frac{\partial}{\partial \theta_i} \hat{\ell}_n^0(\theta) \right) + 2 \sum_{m < n} \text{Cov} \left(\frac{\partial}{\partial \theta_i} \hat{\ell}_m^0(\theta), \frac{\partial}{\partial \theta_i} \hat{\ell}_n^0(\theta) \right) \quad (\text{S155})$$

$$\leq \frac{NG'(\theta)^2}{J} + 2 \sum_{n=1}^{N-1} (N-n)(1 - \epsilon)^{\frac{1}{2} \lfloor \frac{n-1}{c\{1+\log(J)\}} \rfloor} \frac{G'(\theta)^2}{J} \quad (\text{S156})$$

$$\leq \frac{G'(\theta)^2}{J} \left(N + 2 \sum_{n=1}^{N-1} (N-n)(1 - \epsilon)^{\frac{1}{2} \lfloor \frac{n-1}{c\{1+\log(J)\}} \rfloor} \right) \quad (\text{S157})$$

$$\lesssim \frac{G'(\theta)^2 N}{J}. \quad (\text{S158})$$

It then follows that as $\theta \in \Theta \subseteq \mathbb{R}^p$,

$$\text{Var}(\nabla_\theta \hat{\ell}^\alpha(\theta)) \lesssim \frac{Np G'(\theta)^2}{J}. \quad (\text{S159})$$

□

Note that the factor of N that pops up here is due to the use of the unnormalized gradient. If one divides the gradient estimate by $\sqrt{N}$ before usage, the variance does not depend on the horizon. If one divides by N , the error is $O(1/\sqrt{NJ})$.

The variance here of $\frac{Np G'(\theta)^2}{J}$ can be thought of as a lower bound for the variance of MOP- α . We will later see that the variance of MOP- α contains this term, and an additional term that corresponds to the memory of the MOP- α gradient estimate.

Having obtained a variance bound for DMOP-0, we are now in a position to tackle a variance bound for DMOP- α . We analyze the case for $\alpha \in (0, 1)$, using the following proof strategy. Instead

of directly analyzing the variance of the MOP- α gradient estimate, we will analyze the variance of a modified estimator where we truncate the weights at k timesteps, so this estimator only looks at most k timesteps back while accumulating the importance weights. We write

$$\hat{s}_k^\alpha(\theta) := \sum_{n=1}^N \log \left(\frac{\sum_{j=1}^J \prod_{i=n-k}^n \left(\frac{g_{i,j}^{A,F,\theta}}{g_{i,j}^{A,F,\phi}} \right)^{\alpha^{(n-i)}}}{\sum_{j=1}^J \prod_{i=n-k}^{n-1} \left(\frac{g_{i,j}^{A,F,\theta}}{g_{i,j}^{A,F,\phi}} \right)^{\alpha^{(n-i)}}} \right) \quad (\text{S160})$$

for the log-likelihood estimate from the k -truncated estimator with discount factor α . It is the sum of components

$$\hat{s}_{n,k}^\alpha(\theta) := \frac{1}{J} \sum_{j=1}^J \sum_{i=n-k}^n \left(\alpha^{n-i} \log \left(g_{i,j}^{A,F,\theta} \right) - \alpha^{n-i+1} \nabla_\theta \log \left(g_{i-1,j}^{A,F,\theta} \right) \right) \quad (\text{S161})$$

over timesteps N .

This enables us to establish strong mixing for the truncated estimator and its gradient, in order to get a bound on the variance that is $O(NJ^{-1})$. We emphasize that this truncated estimator is only a theoretical construct – we do not actually use the truncated estimator. The truncated estimator, as we note in the discussion, possesses similar theoretical guarantees as the MOP- α estimator, but would require the user to specify the number of timesteps to truncate at instead of a discount factor.

Coupled with a bound ensuring that the k -truncated estimator provides an estimate that is close to that of MOP- α proper, we get a bound on the variance of MOP- α that comprises of the variance of the k -truncated estimator plus the error, for any $k \leq N$. It then holds that the final variance bound on MOP- α is given by the minimum over all k , and is never larger than $O(\psi(\alpha) + NJ^{-1})$ for some function ψ increasing in α .

Theorem S4 (Variance of DMOP- α). *When $\alpha \in (0, 1)$, the variance of the gradient estimate from MOP- α is*

$$\text{Var}(\nabla_\theta \hat{\ell}^\alpha(\theta)) \lesssim \min_{k \leq N} \left(\frac{k^2 G'(\theta)^2 Np}{(1-\alpha)^2 J} + \frac{\alpha^{2k}}{(1-\alpha)^2} Np G'(\theta)^2 \right). \quad (\text{S162})$$

Proof. Using the derivative of the logarithm and that $w_{n,j}^{F,\theta,\alpha} = w_{n,j}^{F,\theta,\alpha,k} = 1$ when $\theta = \phi$,

$$\left\| \nabla_\theta \log \left(\sum_{j=1}^J w_{n,j}^{F,\theta,\alpha} \right) - \nabla_\theta \log \left(\sum_{j=1}^J w_{n,j}^{F,\theta,\alpha,k} \right) \right\|_2 \quad (\text{S163})$$

$$= \left\| \frac{\nabla_\theta \sum_{j=1}^J w_{n,j}^{F,\theta,\alpha}}{\sum_{j=1}^J w_{n,j}^{F,\theta,\alpha}} - \frac{\nabla_\theta \sum_{j=1}^J w_{n,j}^{F,\theta,\alpha,k}}{\sum_{j=1}^J w_{n,j}^{F,\theta,\alpha,k}} \right\|_2 \quad (\text{S164})$$

$$= \left\| \frac{1}{J} \sum_{j=1}^J \nabla_\theta w_{n,j}^{F,\theta,\alpha} - \frac{1}{J} \sum_{j=1}^J \nabla_\theta w_{n,j}^{F,\theta,\alpha,k} \right\|_2 \quad (\text{S165})$$

$$= \left\| \frac{1}{J} \sum_{j=1}^J \frac{\nabla_\theta w_{n,j}^{F,\theta,\alpha}}{w_{n,j}^{F,\theta,\alpha}} - \frac{1}{J} \sum_{j=1}^J \frac{\nabla_\theta w_{n,j}^{F,\theta,\alpha,k}}{w_{n,j}^{F,\theta,\alpha,k}} \right\|_2 \quad (\text{S166})$$

$$= \left\| \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \log(w_{n,j}^{F,\theta,\alpha}) - \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \log(w_{n,j}^{F,\theta,\alpha,k}) \right\|_2. \quad (\text{S167})$$

This lets us bound the cumulative weight discrepancies by

$$\begin{aligned} & \left\| \nabla_{\theta} \log \left(\sum_{j=1}^J w_{n,j}^{F,\theta,\alpha} \right) - \nabla_{\theta} \log \left(\sum_{j=1}^J w_{n,j}^{F,\theta,\alpha,k} \right) \right\|_2 \\ &= \left\| \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \log(w_{n,j}^{F,\theta,\alpha}) - \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \log(w_{n,j}^{F,\theta,\alpha,k}) \right\|_2 \end{aligned} \quad (\text{S168})$$

$$= \left\| \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \left(\log(w_{n,j}^{F,\theta,\alpha}) - \log(w_{n,j}^{F,\theta,\alpha,k}) \right) \right\|_2 \quad (\text{S169})$$

$$= \left\| \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \left(\sum_{i=1}^n \alpha^{(n-i)} \log \left(\frac{g_{i,j}^{A,F,\theta}}{g_{i,j}^{A,F,\phi}} \right) - \sum_{i=n-k}^n \alpha^{(n-i)} \log \left(\frac{g_{i,j}^{A,F,\theta}}{g_{i,j}^{A,F,\phi}} \right) \right) \right\|_2 \quad (\text{S170})$$

$$= \left\| \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \left(\sum_{i=1}^{n-k} \alpha^{(n-i)} \log \left(\frac{g_{i,j}^{A,F,\theta}}{g_{i,j}^{A,F,\phi}} \right) \right) \right\|_2 \quad (\text{S171})$$

$$\leq \frac{1}{J} \sum_{j=1}^J \sum_{i=1}^{n-k} \alpha^{(n-i)} \left\| \nabla_{\theta} \log(g_{i,j}^{A,F,\theta}) \right\|_2 \quad (\text{S172})$$

$$\leq \frac{1}{J} \sum_{j=1}^J \sum_{i=1}^{n-k} \alpha^{(n-i)} p G'(\theta)^2 \quad (\text{S173})$$

$$\leq \sqrt{p} G'(\theta) \frac{\alpha^k - \alpha^n}{1 - \alpha}, \quad (\text{S174})$$

where the second-last line follows from Assumption (A3).

Now, we bound the variance of $\nabla_{\theta} \hat{s}_k^{\alpha}(\theta)$. Recall that we defined

$$\nabla_{\theta} \hat{s}_k^{\alpha}(\theta) = \sum_{n=1}^N \nabla_{\theta} \log \left(\frac{\sum_{j=1}^J \prod_{i=n-k}^n \left(\frac{g_{i,j}^{A,F,\theta}}{g_{i,j}^{A,F,\phi}} \right)^{\alpha^{(n-i)}}}{\sum_{j=1}^J \prod_{i=n-k}^{n-1} \left(\frac{g_{i,j}^{A,F,\theta}}{g_{i,j}^{A,F,\phi}} \right)^{\alpha^{(n-i)}}} \right) \quad (\text{S175})$$

$$= \sum_{n=1}^N \nabla_{\theta} \left(\log \left(\sum_{j=1}^J w_{n,j}^{F,\theta,\alpha,k} \right) - \log \left(\sum_{j=1}^J w_{n-1,j}^{A,F,\theta,\alpha,k} \right) \right). \quad (\text{S176})$$

We can then decompose this expression with the derivative of the logarithm while noting that $w_{n,j}^{F,\theta,\alpha,k} = 1$ whenever $\theta = \phi$, to see that

$$\nabla_{\theta} \hat{s}_k^{\alpha}(\theta) = \sum_{n=1}^N \nabla_{\theta} \left(\log \left(\sum_{j=1}^J w_{n,j}^{F,\theta,\alpha,k} \right) - \log \left(\sum_{j=1}^J w_{n-1,j}^{A,F,\theta,\alpha,k} \right) \right) \quad (\text{S177})$$

$$= \sum_{n=1}^N \left(\frac{\sum_{j=1}^J \nabla_{\theta} w_{n,j}^{F,\theta,\alpha,k}}{\sum_{j=1}^J w_{n,j}^{F,\theta,\alpha,k}} - \frac{\sum_{j=1}^J \nabla_{\theta} w_{n-1,j}^{A,F,\theta,\alpha,k}}{\sum_{j=1}^J w_{n-1,j}^{A,F,\theta,\alpha,k}} \right) \quad (\text{S178})$$

$$= \sum_{n=1}^N \left(\frac{1}{J} \sum_{j=1}^J \nabla_{\theta} w_{n,j}^{F,\theta,\alpha,k} - \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} w_{n-1,j}^{A,F,\theta,\alpha,k} \right) \quad (\text{S179})$$

$$= \frac{1}{J} \sum_{n=1}^N \sum_{j=1}^J \nabla_{\theta} \left(w_{n,j}^{F,\theta,\alpha,k} - w_{n-1,j}^{A,F,\theta,\alpha,k} \right). \quad (\text{S180})$$

It now follows that we need to bound the variance at a single timestep, namely

$$\text{Var} \left(\frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \left(w_{n,j}^{F,\theta,\alpha,k} - w_{n-1,j}^{A,F,\theta,\alpha,k} \right) \right). \quad (\text{S181})$$

We use the derivative of the logarithm yet again to find, noting that $w_{n,j}^{F,\theta,\alpha,k} = 1$ when $\theta = \phi$, that

$$\nabla_{\theta} w_{n,j}^{F,\theta,\alpha,k} = \frac{\nabla_{\theta} w_{n,j}^{F,\theta,\alpha,k}}{w_{n,j}^{F,\theta,\alpha,k}} = \nabla_{\theta} \log(w_{n,j}^{F,\theta,\alpha,k}) = \sum_{i=n-k}^n \alpha^{n-i} \nabla_{\theta} \log \left(\frac{g_{i,j}^{A,F,\theta}}{g_{i,j}^{A,F,\phi}} \right) \quad (\text{S182})$$

$$= \sum_{i=n-k}^n \alpha^{n-i} \nabla_{\theta} \log \left(g_{i,j}^{A,F,\theta} \right). \quad (\text{S183})$$

Then,

$$\text{Var} \left(\nabla_{\theta} (w_{n,j}^{F,\theta,\alpha,k} - w_{n-1,j}^{A,F,\theta,\alpha,k}) \right) \quad (\text{S184})$$

$$= \text{Var} \left(\nabla_{\theta} \left(\sum_{i=n-k}^n \alpha^{n-i} \nabla_{\theta} \log \left(g_{i,j}^{A,F,\theta} \right) - \sum_{i=n-k-1}^{n-1} \alpha^{n-i} \nabla_{\theta} \log \left(g_{i,j}^{A,F,\theta} \right) \right) \right) \quad (\text{S185})$$

$$= \text{Var} \left(\sum_{i=n-k}^n \nabla_{\theta} \left(\alpha^{n-i} \log \left(g_{i,j}^{A,F,\theta} \right) - \alpha^{n-i+1} \log \left(g_{i-1,j}^{A,F,\theta} \right) \right) \right). \quad (\text{S186})$$

Note that

$$\nabla_{\theta} \hat{s}_{n,k}^{\alpha}(\theta) := \frac{1}{J} \sum_{j=1}^J \sum_{i=n-k}^n \nabla_{\theta} \left(\alpha^{n-i} \log \left(g_{i,j}^{A,F,\theta} \right) - \alpha^{n-i+1} \log \left(g_{i-1,j}^{A,F,\theta} \right) \right) \quad (\text{S187})$$

is a function bounded by $C \frac{G'(\theta)(1-\alpha^k)}{1-\alpha} =: CG'(\theta)r$ in each coordinate for some constant C (by Assumption (A3)) that depends only on k timesteps, and not n . Subsequently, we suppress vector and matrix notation for the coordinates of θ and its derivatives, and their variances and covariances, with the variance of a vector interpreted as the sum of the variance of its components. Lemma 2.3 of Karjalainen et al. (2023) gives the required prediction-empirical L^4 bound. For a filtered empirical average, the conditional multinomial-resampling and Bayes-ratio calculation in Equations (S290)–(S296) gives the same order. After centering by Equation (S297), the L^2 error of each of the k

functionals from time $n - k$ to time n is therefore bounded by $Cr\sqrt{p}G'(\theta)J^{-1/2}k$. That is, we have that

$$\mathbb{E} [\|\nabla_{\theta}\hat{s}_{n,k}^{\alpha}(\theta) - \mathbb{E}_{\pi}\nabla_{\theta}\hat{s}_{n,k}^{\alpha}(\theta)\|_2^2]^{1/2} \leq \frac{Cr\sqrt{p}G'(\theta)k}{\sqrt{J}}. \quad (\text{S188})$$

By the bias-variance decomposition, this in turn bounds the variance at a single timestep by $C^2r^2pG'(\theta)^2J^{-1}k^2$. Concretely, we have that

$$\text{Var} \left(\frac{1}{J} \sum_{j=1}^J \sum_{i=n-k}^n \nabla_{\theta} \left(\alpha^{n-i} \log \left(g_{i,j}^{A,F,\theta} \right) - \alpha^{n-i+1} \log \left(g_{i-1,j}^{A,F,\theta} \right) \right) \right) \lesssim \frac{r^2k^2p2G'(\theta)^2}{J}. \quad (\text{S189})$$

It now remains to bound the variance of $\nabla_{\theta}\hat{s}_k^{\alpha}(\theta)$ by considering the covariance of each of the N terms that comprise it. We decompose

$$\begin{aligned} \text{Var}(\nabla_{\theta}\hat{s}_k^{\alpha}(\theta)) &= \text{Var} \left(\sum_{n=1}^N \frac{1}{J} \sum_{j=1}^J \sum_{i=n-k}^n \nabla_{\theta} \left(\alpha^{n-i} \log \left(g_{i,j}^{A,F,\theta} \right) - \alpha^{n-i+1} \log \left(g_{i-1,j}^{A,F,\theta} \right) \right) \right) \\ &= \sum_{n=1}^N \text{Var}(\nabla_{\theta}\hat{s}_{n,k}^{\alpha}(\theta)) + 2 \sum_{m < n} \text{Cov}(\nabla_{\theta}\hat{s}_{m,k}^{\alpha}(\theta), \nabla_{\theta}\hat{s}_{n,k}^{\alpha}(\theta)) \end{aligned} \quad (\text{S190})$$

$$\lesssim \frac{r^2k^2pG'(\theta)^2N}{J} + 2 \sum_{m < n} \text{Cov}(\nabla_{\theta}\hat{s}_{m,k}^{\alpha}(\theta), \nabla_{\theta}\hat{s}_{n,k}^{\alpha}(\theta)). \quad (\text{S191})$$

Similarly to the proof of the MOP-0 case, we use Assumptions (A2) and (A3) that ensure strong mixing. We know from Theorem 3.1 of Karjalainen et al. (2023) that when $\mathbf{M}_{n,n+k}$ is the k -step Markov operator from timestep n and $\beta_{\text{TV}}(M) = \sup_{x,y \in E} \|M(x, \cdot) - M(y, \cdot)\|_{\text{TV}} = \sup_{\mu, \nu \in \mathcal{P}, \mu \neq \nu} \frac{\|\mu M - \nu M\|_{\text{TV}}}{\|\mu - \nu\|_{\text{TV}}}$ is the Dobrushin contraction coefficient of a Markov operator,

$$\beta_{\text{TV}}(\mathbf{M}_{n,n+k}) \leq (1 - \epsilon)^{\lfloor k/(c \log(J)) \rfloor}, \quad (\text{S192})$$

i.e. the mixing time of the particle filter is $O(\log(J))$, where ϵ and c depend on $\bar{M}, \underline{M}, \bar{G}, \underline{G}$ in (A2) and (A3). By Lemma S9, the particle filter itself is strongly mixing, with α -mixing coefficients $a(l) \leq (1 - \epsilon)^{\lfloor l/(c \log(J)) \rfloor}$. Therefore, functions of particles are strongly mixing as well, with α -mixing coefficients bounded by the original (to see this, observe that the σ -algebra of the functionals is contained within the original σ -algebra).

We now derive the α -mixing coefficients of $(\nabla_{\theta}\hat{s}_{n,k}^{\alpha}(\theta))_{n=1}^N$. Observe that $\nabla_{\theta}\hat{s}_{n,k}^{\alpha}$ is strong mixing at lag $k + 1$, as all weights beyond k timesteps are truncated. We therefore have that the mixing time of $\nabla_{\theta}\hat{s}_{n,k}^{\alpha}(\theta)$ is $O(1 + k + \log(J))$, and that the α -mixing coefficients for $(\nabla_{\theta}\hat{s}_{n,k}^{\alpha}(\theta))_{n=1}^N$ are given by $a(l) \leq (1 - \epsilon)^{\lfloor l/(c(1+k+\log(J))) \rfloor}$. Therefore, by Davydov's inequality and the centered L^4 bound just described, we find that

$$\text{Cov}(\nabla_{\theta}\hat{s}_{m,k}^{\alpha}(\theta), \nabla_{\theta}\hat{s}_{n,k}^{\alpha}(\theta)) \lesssim (1 - \epsilon)^{\frac{1}{2} \lfloor \frac{|n-m|}{c(1+k+\log(J))} \rfloor} \frac{r^2pG'(\theta)^2}{J}. \quad (\text{S193})$$

Concretely, Davydov's inequality, applied to the centered variables, gives

$$|\text{Cov}(\nabla_{\theta}\hat{s}_{m,k}^{\alpha}(\theta), \nabla_{\theta}\hat{s}_{n,k}^{\alpha}(\theta))| \quad (\text{S194})$$

$$\leq Ca(n-m)^{1/2} \left\| \nabla_{\theta} \hat{s}_{m,k}^{\alpha}(\theta) - \mathbb{E} \nabla_{\theta} \hat{s}_{m,k}^{\alpha}(\theta) \right\|_{L^4} \left\| \nabla_{\theta} \hat{s}_{n,k}^{\alpha}(\theta) - \mathbb{E} \nabla_{\theta} \hat{s}_{n,k}^{\alpha}(\theta) \right\|_{L^4}. \quad (\text{S195})$$

The predictor-to-filtered calculation cited above and Equation (S297) give, for $q = m, n$,

$$\left\| \nabla_{\theta} \hat{s}_{q,k}^{\alpha}(\theta) - \mathbb{E} \nabla_{\theta} \hat{s}_{q,k}^{\alpha}(\theta) \right\|_{L^4} \lesssim \frac{r\sqrt{p}G'(\theta)}{\sqrt{J}}. \quad (\text{S196})$$

Substitution in Davydov's inequality gives

$$\text{Cov}(\nabla_{\theta} \hat{s}_{m,k}^{\alpha}(\theta), \nabla_{\theta} \hat{s}_{n,k}^{\alpha}(\theta)) \leq (1-\epsilon)^{\frac{1}{2} \lfloor \frac{|n-m|}{c(1+k+\log(J))} \rfloor} \frac{r^2 p G'(\theta)^2}{J}. \quad (\text{S197})$$

Putting it all together, we see that

$$\text{Var}(\nabla_{\theta} \hat{s}_k(\theta)) = \sum_{n=1}^N \text{Var}(\nabla_{\theta} \hat{s}_{n,k}^{\alpha}(\theta)) + 2 \sum_{m < n} \text{Cov}(\nabla_{\theta} \hat{s}_{m,k}^{\alpha}(\theta), \nabla_{\theta} \hat{s}_{n,k}^{\alpha}(\theta)) \quad (\text{S198})$$

$$\leq \frac{Nr^2 p k^2 G'(\theta)^2}{J} + 2 \sum_{n=1}^{N-1} (N-n)(1-\epsilon)^{\frac{1}{2} \lfloor \frac{n}{c(1+k+\log(J))} \rfloor} \frac{r^2 p G'(\theta)^2}{J} \quad (\text{S199})$$

$$\leq \frac{r^2 p k^2 G'(\theta)^2}{J} \left(N + 2 \sum_{n=1}^{N-1} (N-n)(1-\epsilon)^{\frac{1}{2} \lfloor \frac{n}{c(1+k+\log(J))} \rfloor} \right) \quad (\text{S200})$$

$$\lesssim \frac{k^2 r^2 p G'(\theta)^2 N}{J}. \quad (\text{S201})$$

To continue, define the per-time truncation discrepancy

$$\Delta_{n,k}(\theta) := \nabla_{\theta} \hat{\ell}_n^{\alpha}(\theta) - \nabla_{\theta} \hat{s}_{n,k}^{\alpha}(\theta),$$

so that

$$\nabla_{\theta} \hat{\ell}^{\alpha}(\theta) - \nabla_{\theta} \hat{s}_k^{\alpha}(\theta) = \sum_{n=1}^N \Delta_{n,k}(\theta).$$

The pointwise calculation above shows that each coordinate of the tail satisfies

$$\|\Delta_{n,k}(\theta)\|_2 \lesssim \frac{\alpha^k}{1-\alpha} \sqrt{p} G'(\theta).$$

However, we do not use the deterministic bound

$$\left\| \sum_{n=1}^N \Delta_{n,k}(\theta) \right\|_2 \leq \sum_{n=1}^N \|\Delta_{n,k}(\theta)\|_2,$$

since this would introduce an unnecessary factor of N^2 after squaring.

Instead, we bound the variance of the cumulative truncation error. By Assumptions (A2)–(A3) and the forgetting bound of Karjalainen et al. (2023), the underlying particle process is strongly mixing. Since $\Delta_{n,k}$ is a geometrically weighted tail functional of bounded score increments, the same covariance-summation argument used for the truncated estimator gives

$$|\text{Cov}(\Delta_{m,k}(\theta), \Delta_{n,k}(\theta))| \lesssim \frac{\alpha^{2k}}{(1-\alpha)^2} p G'(\theta)^2 \rho(|n-m|)$$

for some summable sequence $\rho(\cdot)$, with

$$\sum_{\ell \geq 0} \rho(\ell) \lesssim 1.$$

Consequently,

$$\begin{aligned} \text{Var}\left(\nabla_{\theta} \hat{\ell}^{\alpha}(\theta) - \nabla_{\theta} \hat{s}_k^{\alpha}(\theta)\right) &= \text{Var}\left(\sum_{n=1}^N \Delta_{n,k}(\theta)\right) \\ &= \sum_{n=1}^N \text{Var}(\Delta_{n,k}(\theta)) + 2 \sum_{m < n} \text{Cov}(\Delta_{m,k}(\theta), \Delta_{n,k}(\theta)) \\ &\lesssim N \frac{\alpha^{2k}}{(1-\alpha)^2} pG'(\theta)^2. \end{aligned}$$

We now compare the variance of the full estimator to the variance of the truncated estimator. For any random vectors X, Y ,

$$\text{Var}(Y) \leq 2\text{Var}(X) + 2\text{Var}(Y - X).$$

Applying this with

$$Y = \nabla_{\theta} \hat{\ell}^{\alpha}(\theta), \quad X = \nabla_{\theta} \hat{s}_k^{\alpha}(\theta),$$

and using the covariance-summation bound above yields

$$\begin{aligned} \text{Var}\left(\nabla_{\theta} \hat{\ell}^{\alpha}(\theta)\right) &\leq 2\text{Var}(\nabla_{\theta} \hat{s}_k^{\alpha}(\theta)) + 2\text{Var}\left(\nabla_{\theta} \hat{\ell}^{\alpha}(\theta) - \nabla_{\theta} \hat{s}_k^{\alpha}(\theta)\right) \\ &\lesssim \frac{k^2 N pG'(\theta)^2}{(1-\alpha)^2 J} + \frac{\alpha^{2k}}{(1-\alpha)^2} N pG'(\theta)^2. \end{aligned}$$

Since this holds for every $k \leq N$, the result follows by minimizing over k . □

Theorem S5 (MSE and variance bounds for DMOP- α). *Suppose the observations are fixed, $J \geq 2$, $\alpha \in [0, 1]$, $\theta = \phi$, Assumptions (A2), (A3), and (A5) hold, and multinomial resampling is used. There are constants $C < \infty$ and $\varrho \in (0, 1)$, independent of N, J , and α , such that*

$$\begin{aligned} \mathbb{E} \left\| \nabla_{\theta} \ell(\theta) - \nabla_{\theta} \hat{\ell}^{\alpha}(\theta) \right\|_2^2 &\lesssim pG'(\theta)^2 \left[\frac{1}{J} \sum_{n=1}^N \left(\sum_{r=0}^{n-1} \alpha^r \right)^2 + \frac{1}{J^2} \left(\sum_{n=1}^N \sum_{r=0}^{n-1} \alpha^r \right)^2 + N^2 \left\{ \frac{\varrho(1-\alpha)}{(1-\varrho)(1-\alpha\varrho)} \right\}^2 \right], \end{aligned} \quad (\text{S202})$$

$$\text{Var}\left(\nabla_{\theta} \hat{\ell}^{\alpha}(\theta)\right) \lesssim pG'(\theta)^2 \min \left\{ N^2, \frac{1}{J} \sum_{n=1}^N \left(\sum_{r=0}^{n-1} \alpha^r \right)^2 \right\}. \quad (\text{S203})$$

For $0 \leq \alpha < 1$, Equation (S202) implies

$$\mathbb{E} \left\| \nabla_{\theta} \ell(\theta) - \nabla_{\theta} \hat{\ell}^{\alpha}(\theta) \right\|_2^2 \lesssim pG'(\theta)^2 \left\{ \frac{N}{J(1-\alpha)^2} + N^2(1-\alpha)^2 + \frac{N^2}{J^2(1-\alpha)^2} \right\}. \quad (\text{S204})$$

Proof. Theorem S6, proved in Section S7, gives Equation (S202). Lemma S13 gives Equation (S203). Finally, Equation (S321) gives Equation (S204). □

S7 Sharper Bias-Variance Analysis

We now establish the estimates used in Theorem S5. Write $\nabla_\theta \ell_n^\alpha(\theta)$ for the n th conditional score produced by the deterministic differentiated population recursion. At $\theta = \phi$, its $\alpha = 1$ version is the exact conditional score $\nabla_\theta \ell_n(\theta)$. A differentiated particle is the particle carried by the differentiated algorithm: it includes its location and the corresponding simulator and weight tangents. Thus the particle-state bounds and mixing statements below apply directly to the differentiable particles used by DMOP; there is no separate “physical particle” whose tangent must be treated afterward. Throughout this section, the observations are fixed. Expectations and variances are over the simulator and particle-filter random variables, conditional on those observations.

We use Assumptions (A3) and (A5) in the precise pathwise sense already required by the differentiated MOP construction. For fixed simulator random numbers and fixed resampling indices, Assumption (A5) says that the recursively simulated state path and the composed log measurement densities are continuously differentiable in θ near ϕ , and that the joint law of the base simulator random numbers does not depend on θ . The derivative in the last clause of Assumption (A3) is the resulting total derivative: uniformly in n, j and in a neighborhood of ϕ , after enlarging $G'(\theta)$ to bound that neighborhood if necessary,

$$\left\| \nabla_\theta \log g_{n,j}^{A,F,\theta} \right\|_\infty \leq G'(\theta). \quad (\text{S205})$$

These two existing assumptions also justify differentiation under the finite-horizon simulator expectations used below. Indeed, the product rule and Equation (S205) give, pathwise,

$$\begin{aligned} & \left\| \nabla_\theta \prod_{s=1}^t G_s(X_s) \right\|_\infty \\ &= \left\| \prod_{s=1}^t G_s(X_s) \sum_{s=1}^t \nabla_\theta \log G_s(X_s) \right\|_\infty \\ &\leq \prod_{s=1}^t G_s(X_s) \sum_{s=1}^t \left\| \nabla_\theta \log G_s(X_s) \right\|_\infty \\ &\leq t \bar{G}^t G'(\theta). \end{aligned} \quad (\text{S206})$$

The final expression is deterministic and finite for each fixed t . For each coordinate of θ , the mean-value theorem therefore bounds every difference quotient about ϕ by the same quantity $t \bar{G}^t G'(\theta)$. The dominated convergence theorem gives

$$\nabla_\theta \mathbb{E}_U \left[\prod_{s=1}^t G_s(X_s) \right] = \mathbb{E}_U \left[\nabla_\theta \prod_{s=1}^t G_s(X_s) \right]. \quad (\text{S207})$$

Thus the required interchange follows from the clarified Assumptions (A3) and (A5); it is not an additional assumption. Geometric population forgetting follows from the backward-kernel contraction under Assumption (A2), and the particle weak-error bound follows from the current-particle telescoping argument in Lemma S11.

For the path functionals below, evaluated at $\theta = \phi$, let ρ_t be the posterior path measure of the differentiable particle. Thus, for every integrable path functional h ,

$$\rho_t(h) = \frac{\mathbb{E} [h(X_{0:t}) \prod_{s=1}^t G_s(X_s)]}{\mathbb{E} [\prod_{s=1}^t G_s(X_s)]}. \quad (\text{S208})$$

Here X_s already includes its tangent, as stipulated above, while $G_s(X_s)$ is the existing measurement potential. Equation (S208) is therefore the usual posterior path measure on the state used by the differentiated particle filter, not an additional extension of it.

Unrolling the existing filtered-weight recursion at $\theta = \phi$ gives

$$\nabla_{\theta} \log w_{n,j}^{F,\theta} \Big|_{\theta=\phi} = \sum_{r=0}^{n-1} \alpha^r \nabla_{\theta} \log g_{n-r,j}^{A,F,\theta} \Big|_{\theta=\phi}. \quad (\text{S209})$$

For a complete ancestral path, write h_k , only in this section, for the measurement-score contribution

$$h_k = \nabla_{\theta} \log g_k^{A,F,\theta} \Big|_{\theta=\phi} \quad (\text{S210})$$

carried on that path. In the path-space notation of Chan et al. (2018), $m(\xi_t)(h_k)$ is its empirical average among the extended ancestral paths at time t , and Equation (S208) gives its population counterpart.

We now identify the lag coefficients in the conditional score. At $\theta = \phi$, all weights equal one, so differentiating the two sample averages in the definition of $L_n^{A,\theta,\alpha}$ gives

$$\begin{aligned} \nabla_{\theta} \hat{\ell}_n^{\alpha}(\theta) &= \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \log w_{n,j}^{F,\theta} \Big|_{\theta=\phi} - \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \log w_{n,j}^{P,\theta} \Big|_{\theta=\phi} \\ &= \frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \log w_{n,j}^{F,\theta} \Big|_{\theta=\phi} - \frac{\alpha}{J} \sum_{j=1}^J \nabla_{\theta} \log w_{n-1,j}^{F,\theta} \Big|_{\theta=\phi}, \end{aligned} \quad (\text{S211})$$

where the second equality uses $w_{n,j}^{P,\theta} = (w_{n-1,j}^{F,\theta})^{\alpha}$. In the MOP- α recursion, the discounted-weight accumulation and process propagation preserve the index j and occur before the time- n resampling step. The resampling index k_j enters only afterward, through

$$w_{n,j}^{F,\theta} = w_{n,k_j}^{P,\theta} \frac{g_{n,k_j}^{\theta}}{g_{n,k_j}^{\phi}}. \quad (\text{S212})$$

Consequently, the second average in Equation (S211) is pathwise

$$\frac{1}{J} \sum_{j=1}^J \nabla_{\theta} \log w_{n,j}^{P,\theta} \Big|_{\theta=\phi} = \frac{\alpha}{J} \sum_{j=1}^J \nabla_{\theta} \log w_{n-1,j}^{F,\theta} \Big|_{\theta=\phi}, \quad (\text{S213})$$

which is the empirical average over the time- $(n-1)$ filtered genealogy. Substituting Equation (S209) into Equation (S211) and changing the index in the second sum give

$$\begin{aligned} \nabla_{\theta} \hat{\ell}_n^{\alpha}(\theta) &= \sum_{r=0}^{n-1} \alpha^r m(\xi_n)(h_{n-r}) - \sum_{r=1}^{n-1} \alpha^r m(\xi_{n-1})(h_{n-r}) \\ &= m(\xi_n)(h_n) + \sum_{r=1}^{n-1} \alpha^r \{m(\xi_n)(h_{n-r}) - m(\xi_{n-1})(h_{n-r})\}. \end{aligned} \quad (\text{S214})$$

Consequently, define

$$\begin{aligned}\Delta_{n,0}^J &= m(\xi_n)(h_n), & \Delta_{n,0} &= \rho_n(h_n), \\ \Delta_{n,r}^J &= m(\xi_n)(h_{n-r}) - m(\xi_{n-1})(h_{n-r}), & \Delta_{n,r} &= \rho_n(h_{n-r}) - \rho_{n-1}(h_{n-r}), \quad 1 \leq r < n.\end{aligned}\quad (\text{S215})$$

We now verify that the right-hand population quantities are the lag coefficients of the deterministic differentiated recursion. Equations (S207) and (S210) give, at $\theta = \phi$,

$$\begin{aligned}\nabla_\theta \log \mathbb{E}_U \left[\prod_{s=1}^n G_s(X_s) \right] &= \frac{\mathbb{E}_U [\prod_{s=1}^n G_s(X_s) \sum_{s=1}^n h_s]}{\mathbb{E}_U [\prod_{s=1}^n G_s(X_s)]} \\ &= \sum_{s=1}^n \rho_n(h_s).\end{aligned}\quad (\text{S216})$$

Subtracting the same identity at time $n-1$ gives

$$\begin{aligned}\nabla_\theta \ell_n(\theta) &= \sum_{s=1}^n \rho_n(h_s) - \sum_{s=1}^{n-1} \rho_{n-1}(h_s) \\ &= \rho_n(h_n) + \sum_{r=1}^{n-1} \{\rho_n(h_{n-r}) - \rho_{n-1}(h_{n-r})\} \\ &= \sum_{r=0}^{n-1} \Delta_{n,r}.\end{aligned}\quad (\text{S217})$$

Finally, differentiating the two population weighted averages in the deterministic MOP- α recursion and using Equation (S207) gives

$$\begin{aligned}\nabla_\theta \ell_n^\alpha(\theta) &= \sum_{r=0}^{n-1} \alpha^r \rho_n(h_{n-r}) - \sum_{r=1}^{n-1} \alpha^r \rho_{n-1}(h_{n-r}) \\ &= \rho_n(h_n) + \sum_{r=1}^{n-1} \alpha^r \{\rho_n(h_{n-r}) - \rho_{n-1}(h_{n-r})\} \\ &= \sum_{r=0}^{n-1} \alpha^r \Delta_{n,r}.\end{aligned}\quad (\text{S218})$$

Lemma S10 (Lagwise first-derivative forgetting). *Suppose $\theta = \phi$. Under Assumptions (A2), (A3), and (A5), there are constants $C < \infty$ and $\varrho \in (0, 1)$ such that, for $0 \leq r < n$,*

$$\|\Delta_{n,r}\|_2 \leq C\sqrt{p} G'(\theta) \varrho^r. \quad (\text{S219})$$

Proof. We derive the result from the posterior update and the contraction of the posterior backward kernels.

First suppose that $1 \leq r < n$, and put $k = n - r$. Fix a coordinate $i \in \{1, \dots, p\}$, and let $h_{k,i}$ be the corresponding coordinate of h_k in Equation (S210). By Equation (S205),

$$|h_{k,i}| \leq G'(\theta). \quad (\text{S220})$$

Because $h_{k,i}$ is measurable with respect to the canonical extended path through time $k \leq n - 1$, Equation (S208) and conditional expectation with respect to the new base simulator variable give

$$\begin{aligned} \rho_n(h_{k,i}) &= \frac{\mathbb{E}_U \left[h_{k,i} \prod_{s=1}^{n-1} G_s(X_s) \mathbb{E}_U \{G_n(X_n) \mid X_{n-1}\} \right]}{\mathbb{E}_U \left[\prod_{s=1}^{n-1} G_s(X_s) \mathbb{E}_U \{G_n(X_n) \mid X_{n-1}\} \right]} \\ &= \frac{\rho_{n-1}\{h_{k,i} M_n(G_n)\}}{\rho_{n-1} M_n(G_n)}, \end{aligned} \quad (\text{S221})$$

where the first equality uses the independence of the new base simulator variable from the path through time $n - 1$. The second equality first uses the identity below and then divides its numerator and denominator by $\mathbb{E}_U \{\prod_{s=1}^{n-1} G_s(X_s)\}$. Using the manuscript's kernel notation, the required identity is

$$M_n(G_n)(x_{n-1}) = \int M_n(x_{n-1}, dx_n) G_n(x_n). \quad (\text{S222})$$

The bounds on G_n and the fact that M_n is a Markov kernel imply for every x_{n-1} that

$$M_n(G_n)(x_{n-1}) \geq \underline{G} \int M_n(x_{n-1}, dx_n) = \underline{G}, \quad (\text{S223})$$

$$M_n(G_n)(x_{n-1}) \leq \bar{G} \int M_n(x_{n-1}, dx_n) = \bar{G}. \quad (\text{S224})$$

We next control how much information about the path through time k can remain at time $n - 1$. For $s \geq 1$, the posterior backward kernel from X_s to X_{s-1} , written only in this proof as B_s , is

$$B_s(x_s, dx_{s-1}) = \frac{\pi_{s-1}(dx_{s-1}) f_{X_s|X_{s-1}}(x_s \mid x_{s-1}; \theta)}{\int \pi_{s-1}(du) f_{X_s|X_{s-1}}(x_s \mid u; \theta)}. \quad (\text{S225})$$

Assumption (A2) gives the following minorization for every measurable set A :

$$\begin{aligned} B_s(x_s, A) &= \frac{\int_A \pi_{s-1}(dx_{s-1}) f_{X_s|X_{s-1}}(x_s \mid x_{s-1}; \theta)}{\int \pi_{s-1}(du) f_{X_s|X_{s-1}}(x_s \mid u; \theta)} \\ &\geq \frac{\underline{M} \pi_{s-1}(A)}{\bar{M} \int \pi_{s-1}(du)} \\ &= \frac{\underline{M}}{\bar{M}} \pi_{s-1}(A). \end{aligned} \quad (\text{S226})$$

Put $\epsilon = \underline{M}/\bar{M}$. If $\epsilon < 1$, Equation (S226) permits the decomposition

$$B_s(x_s, \cdot) = \epsilon \pi_{s-1}(\cdot) + (1 - \epsilon) R_s(x_s, \cdot), \quad (\text{S227})$$

where

$$R_s(x_s, A) = \frac{B_s(x_s, A) - \epsilon \pi_{s-1}(A)}{1 - \epsilon} \quad (\text{S228})$$

is a Markov kernel. Hence, for every bounded real function f and any x_s, x'_s ,

$$\begin{aligned} B_s(f)(x_s) - B_s(f)(x'_s) &= (1 - \epsilon) \{R_s(f)(x_s) - R_s(f)(x'_s)\}, \\ |B_s(f)(x_s) - B_s(f)(x'_s)| &\leq (1 - \epsilon) \text{osc}(f). \end{aligned} \quad (\text{S229})$$

If $\epsilon = 1$, Equation (S226) implies $B_s(x_s, \cdot) = \pi_{s-1}(\cdot)$, so the left-hand side is zero. Taking the supremum over x_s, x'_s gives

$$\text{osc}\{B_s(f)\} \leq (1 - \epsilon) \text{osc}(f). \quad (\text{S230})$$

The function $h_{k,i}$ is a bounded function of the differentiable particle X_k . Repeated conditioning through the backward kernels therefore gives

$$\mathbb{E}_{\rho_{n-1}}\{h_{k,i}(X_k) \mid X_{n-1} = x_{n-1}\} = (B_{n-1}B_{n-2} \cdots B_{k+1}h_{k,i})(x_{n-1}). \quad (\text{S231})$$

There are $n - 1 - k = r - 1$ backward kernels in this display. Moreover, Equation (S220) gives

$$\begin{aligned} \text{osc}(h_{k,i}) &\leq 2\|h_{k,i}\|_\infty \\ &\leq 2G'(\theta). \end{aligned} \quad (\text{S232})$$

If $r = 1$, Equation (S231) contains no backward kernel. If $r \geq 2$, applying Equation (S230) successively gives

$$\begin{aligned} &\text{osc}[\mathbb{E}_{\rho_{n-1}}\{h_{k,i} \mid X_{n-1} = \cdot\}] \\ &\leq 2G'(\theta) \begin{cases} 1, & r = 1, \\ (1 - \epsilon)^{r-1}, & r \geq 2. \end{cases} \end{aligned} \quad (\text{S233})$$

To finish, abbreviate the conditional expectation on the left of Equation (S233) by H , only for the next calculation. Since the X_{n-1} -marginal of ρ_{n-1} is π_{n-1} , the tower property gives

$$\rho_{n-1}(h_{k,i}) = \pi_{n-1}(H), \quad (\text{S234})$$

$$\rho_{n-1}\{h_{k,i}M_n(G_n)\} = \pi_{n-1}\{HM_n(G_n)\}. \quad (\text{S235})$$

Substituting Equations (S234)–(S235) into Equation (S221) and then subtracting $\rho_{n-1}(h_{k,i})$ gives the exact covariance identity

$$\begin{aligned} \Delta_{n,r,i} &= \rho_n(h_{k,i}) - \rho_{n-1}(h_{k,i}) \\ &= \frac{\pi_{n-1}\{HM_n(G_n)\} - \pi_{n-1}(H)\pi_{n-1}M_n(G_n)}{\pi_{n-1}M_n(G_n)} \\ &= \frac{\pi_{n-1}[\{H - \pi_{n-1}(H)\}\{M_n(G_n) - \pi_{n-1}M_n(G_n)\}]}{\pi_{n-1}M_n(G_n)}. \end{aligned} \quad (\text{S236})$$

For any probability measure μ and bounded real f ,

$$|f(x) - \mu(f)| \leq \text{osc}(f). \quad (\text{S237})$$

Equations (S223)– (S224) also imply

$$\text{osc}\{M_n(G_n)\} \leq \bar{G} - \underline{G}. \quad (\text{S238})$$

Taking absolute values in Equation (S236), using Equations (S237)– (S238), and then using Equation (S233), yields

$$\begin{aligned} |\Delta_{n,r,i}| &\leq \frac{\text{osc}(H) \text{osc}\{M_n(G_n)\}}{\underline{G}} \\ &\leq \frac{2(\bar{G} - \underline{G})}{\underline{G}} G'(\theta) \begin{cases} 1, & r = 1, \\ (1 - \epsilon)^{r-1}, & r \geq 2. \end{cases} \end{aligned} \quad (\text{S239})$$

Choose any $\varrho \in (1 - \epsilon, 1)$. Since $r \geq 1$,

$$\begin{cases} 1, & r = 1, \\ (1 - \epsilon)^{r-1}, & r \geq 2 \end{cases} \leq \varrho^{r-1} = \varrho^{-1} \varrho^r. \quad (\text{S240})$$

Absorbing the fixed factor ϱ^{-1} into C , summing the squared coordinate bounds, and taking square roots give

$$\begin{aligned} \|\Delta_{n,r}\|_2 &= \left\{ \sum_{i=1}^p |\Delta_{n,r,i}|^2 \right\}^{1/2} \\ &\leq C\sqrt{p} G'(\theta) \varrho^r, \quad 1 \leq r < n. \end{aligned} \quad (\text{S241})$$

Finally, when $r = 0$, Equation (S215) and Equation (S205) give

$$\begin{aligned} \|\Delta_{n,0}\|_2 &= \|\rho_n(h_n)\|_2 \\ &\leq \rho_n(\|h_n\|_2) \\ &\leq \sqrt{p} G'(\theta) \\ &\leq C\sqrt{p} G'(\theta) \varrho^0. \end{aligned} \quad (\text{S242})$$

Equations (S241) and (S242) prove Equation (S219). □

Lemma S11 (Uniform weak bias of current-particle averages). *Suppose $\theta = \phi$. Under Assumptions (A2), (A3), and (A5), with multinomial resampling, there is $C < \infty$ such that, for every bounded real function f of the current differentiable particle and every $t \geq 1$,*

$$|\mathbb{E}m(\xi_t)(f) - \rho_t(f)| \leq \frac{C}{J} \text{osc}(f). \quad (\text{S243})$$

The constant is independent of t and J .

Proof. The predictor cloud consists of the differentiable particles themselves: their locations, simulator tangents, and weight tangents are coordinates of the particle state, rather than an external process appended to the cloud. Thus f is a current-state functional in the Feynman–Kac model

already defined in Section S3. The current-state marginal of ρ_t is the filtering distribution, and hence

$$\rho_t(f) = \pi_t(f) = \Psi_t(\eta_t)(f). \quad (\text{S244})$$

Conditionally on the prediction cloud η_t^J , multinomial resampling draws the filtering particles independently from $\Psi_t(\eta_t^J)$. Therefore

$$\begin{aligned} \mathbb{E}\{m(\xi_t)(f) \mid \eta_t^J\} &= \frac{1}{J} \sum_{j=1}^J \mathbb{E}\{f(x_{t,j}^F) \mid \eta_t^J\} \\ &= \frac{1}{J} \sum_{j=1}^J \Psi_t(\eta_t^J)(f) \\ &= \Psi_t(\eta_t^J)(f). \end{aligned} \quad (\text{S245})$$

Take $\Phi_{t,t}$ to be the identity. For $1 \leq p \leq t$, define

$$\begin{aligned} Q_{t,t}(f)(x) &= G_t(x)f(x), \\ Q_{p,t}(f)(x) &= G_p(x)M_{p+1}\{Q_{p+1,t}(f)\}(x), \quad p < t. \end{aligned} \quad (\text{S246})$$

The definitions of $\Phi_{p,t}$ and Ψ_t give, for every probability measure μ ,

$$\Psi_t\{\Phi_{p,t}(\mu)\}(f) = \frac{\mu Q_{p,t}(f)}{\mu Q_{p,t}(1)}. \quad (\text{S247})$$

For $p < t$,

$$Q_{p,t}(1)(x) = G_p(x) \int f_{X_{p+1}|X_p}(u \mid x; \theta) Q_{p+1,t}(1)(u) du. \quad (\text{S248})$$

Thus, for every x, x' ,

$$\begin{aligned} \frac{Q_{p,t}(1)(x)}{Q_{p,t}(1)(x')} &\leq \frac{\bar{G}\bar{M} \int Q_{p+1,t}(1)(u) du}{\underline{G}\underline{M} \int Q_{p+1,t}(1)(u) du} \\ &= \frac{\bar{G}\bar{M}}{\underline{G}\underline{M}}. \end{aligned} \quad (\text{S249})$$

For $p = t$, the same ratio is at most $\bar{G}/\underline{G}$. Put

$$R = \frac{\bar{G}\bar{M}}{\underline{G}\underline{M}} \geq 1. \quad (\text{S250})$$

Then

$$\sup_{x, x'} \frac{Q_{p,t}(1)(x)}{Q_{p,t}(1)(x')} \leq R, \quad 1 \leq p \leq t. \quad (\text{S251})$$

For $p \leq s < t$, the normalized forward kernel is

$$P_{s,s+1}^{(t)}(x, A) = \frac{\int_A f_{X_{s+1}|X_s}(u \mid x; \theta) Q_{s+1,t}(1)(u) du}{\int f_{X_{s+1}|X_s}(u \mid x; \theta) Q_{s+1,t}(1)(u) du}. \quad (\text{S252})$$

Define, only within this proof,

$$\nu_{s,t}(A) = \frac{\int_A Q_{s+1,t}(1)(u) du}{\int Q_{s+1,t}(1)(u) du}. \quad (\text{S253})$$

Assumption (A2) gives

$$\begin{aligned} P_{s,s+1}^{(t)}(x, A) &\geq \frac{\underline{M} \int_A Q_{s+1,t}(1)(u) du}{\bar{M} \int Q_{s+1,t}(1)(u) du} \\ &= \frac{\underline{M}}{\bar{M}} \nu_{s,t}(A). \end{aligned} \quad (\text{S254})$$

Consequently, with

$$\varrho = 1 - \frac{\underline{M}}{2\bar{M}} \in (0, 1), \quad (\text{S255})$$

the Dobrushin coefficient of every kernel in Equation (S252) is at most ϱ . The factor $G_s(x)$ cancels from the normalized conditional expectation, and hence

$$\frac{Q_{p,t}(f)}{Q_{p,t}(1)} = P_{p,p+1}^{(t)} P_{p+1,p+2}^{(t)} \cdots P_{t-1,t}^{(t)} f. \quad (\text{S256})$$

It follows that

$$\text{osc} \left\{ \frac{Q_{p,t}(f)}{Q_{p,t}(1)} \right\} \leq \varrho^{t-p} \text{osc}(f). \quad (\text{S257})$$

We next bound one complete selection–mutation update. For $p = 1$, put $\mu = \eta_1 = \pi_0 M_1$. For $p \geq 2$, condition on η_{p-1}^J and put

$$\mu = \Phi_p(\eta_{p-1}^J). \quad (\text{S258})$$

In either case, the particles forming η_p^J are conditionally independent with common law μ . Within this calculation, put

$$q = \frac{Q_{p,t}(1)}{\mu Q_{p,t}(1)}, \quad a = \frac{\mu Q_{p,t}(f)}{\mu Q_{p,t}(1)}, \quad d = \frac{Q_{p,t}(f - a)}{\mu Q_{p,t}(1)}. \quad (\text{S259})$$

Then

$$\mu(q) = 1, \quad \mu(d) = 0, \quad (\text{S260})$$

and Equation (S247) gives

$$\begin{aligned} \Psi_t\{\Phi_{p,t}(\eta_p^J)\}(f) &- \Psi_t\{\Phi_{p,t}(\mu)\}(f) \\ &= \frac{\eta_p^J Q_{p,t}(f)}{\eta_p^J Q_{p,t}(1)} - a \\ &= \frac{\eta_p^J Q_{p,t}(f - a)}{\eta_p^J Q_{p,t}(1)} \end{aligned}$$

$$= \frac{\eta_p^J(d)}{\eta_p^J(q)}. \quad (\text{S261})$$

Equation (S251) implies

$$R^{-1} \leq q(x) \leq R, \quad \eta_p^J(q) \geq R^{-1}, \quad \|q - 1\|_\infty \leq R. \quad (\text{S262})$$

Writing

$$F = \frac{Q_{p,t}(f)}{Q_{p,t}(1)}, \quad (\text{S263})$$

we have $a = \mu(qF)$ and $\mu(q) = 1$. Thus a is an average of F under the probability measure $q d\mu$, so

$$\inf F \leq a \leq \sup F. \quad (\text{S264})$$

Since $d = q(F - a)$, Equations (S262), (S264), and (S257) give

$$\begin{aligned} \|d\|_\infty &\leq R \operatorname{osc}(F) \\ &\leq R \varrho^{t-p} \operatorname{osc}(f). \end{aligned} \quad (\text{S265})$$

Conditional independence and $\mu(d) = 0$ give

$$\begin{aligned} \mathbb{E}\{\eta_p^J(d) \mid \mu\} &= 0, \\ \mathbb{E}\{\eta_p^J(d)^2 \mid \mu\} &= \frac{1}{J} \mu(d^2), \\ \mathbb{E}\{(\eta_p^J(q) - 1)^2 \mid \mu\} &= \frac{1}{J} \mu\{(q - 1)^2\}. \end{aligned} \quad (\text{S266})$$

Consequently,

$$\begin{aligned} &\mathbb{E} \left\{ \frac{\eta_p^J(d)}{\eta_p^J(q)} \middle| \mu \right\} \\ &= -\mathbb{E} \left\{ \frac{\eta_p^J(d) \{ \eta_p^J(q) - 1 \}}{\eta_p^J(q)} \middle| \mu \right\}. \end{aligned} \quad (\text{S267})$$

Conditional Cauchy–Schwarz and Equations (S262)–(S266) yield

$$\begin{aligned} &\left| \mathbb{E} \left\{ \frac{\eta_p^J(d)}{\eta_p^J(q)} \middle| \mu \right\} \right| \\ &\leq R [\mathbb{E}\{\eta_p^J(d)^2 \mid \mu\} \mathbb{E}\{(\eta_p^J(q) - 1)^2 \mid \mu\}]^{1/2} \\ &= \frac{R}{J} [\mu(d^2) \mu\{(q - 1)^2\}]^{1/2} \\ &\leq \frac{C}{J} \varrho^{t-p} \operatorname{osc}(f). \end{aligned} \quad (\text{S268})$$

It remains to write the exact full-update telescope. Since $\eta_t = \Phi_{1,t}(\eta_1)$,

$$\mathbb{E} \Psi_t(\eta_t^J)(f) - \Psi_t(\eta_t)(f)$$

$$\begin{aligned}
&= \mathbb{E} [\Psi_t\{\Phi_{1,t}(\eta_1^J)\}(f) - \Psi_t\{\Phi_{1,t}(\eta_1)\}(f)] \\
&\quad + \sum_{p=2}^t \mathbb{E} [\Psi_t\{\Phi_{p,t}(\eta_p^J)\}(f) - \Psi_t\{\Phi_{p,t}(\Phi_p(\eta_{p-1}^J))\}(f)].
\end{aligned} \tag{S269}$$

Indeed, for $2 \leq p \leq t$,

$$\Phi_{p,t} \circ \Phi_p = \Phi_{p-1,t}, \tag{S270}$$

so consecutive terms cancel. The only terms left are $\Psi_t(\eta_t^J)(f)$ and $-\Psi_t(\eta_t)(f)$.

Apply Equation (S268) to every summand in Equation (S269). Then

$$\begin{aligned}
|\mathbb{E}\Psi_t(\eta_t^J)(f) - \Psi_t(\eta_t)(f)| &\leq \frac{C \operatorname{osc}(f)}{J} \sum_{p=1}^t \varrho^{t-p} \\
&\leq \frac{C}{J(1-\varrho)} \operatorname{osc}(f).
\end{aligned} \tag{S271}$$

Absorb $(1-\varrho)^{-1}$ into C and combine this display with Equations (S244) and (S245). $\square$

Lemma S12 (Particle bias of the full DMOP- α score). *Suppose $\theta = \phi$. Under the conditions of Lemma S11, for every $\alpha \in [0, 1]$,*

$$\left\| \mathbb{E} \nabla_{\theta} \hat{\ell}^{\alpha}(\theta) - \nabla_{\theta} \ell^{\alpha}(\theta) \right\|_2 \lesssim \frac{\sqrt{p} G'(\theta)}{J} \sum_{n=1}^N \sum_{r=0}^{n-1} \alpha^r, \tag{S272}$$

$$\left\| \mathbb{E} \nabla_{\theta} \hat{\ell}^{\alpha}(\theta) - \nabla_{\theta} \ell^{\alpha}(\theta) \right\|_2 \leq \frac{C \sqrt{p} G'(\theta) N^2}{J}. \tag{S273}$$

The constants are independent of N, J , and α .

Proof. We first record the summation-by-parts identity used again in the variance proof. Sum Equation (S214) over $n = 1, \dots, N$ and put $k = n - r$. Since the sums are finite,

$$\nabla_{\theta} \hat{\ell}^{\alpha}(\theta) = \sum_{k=1}^N \left[m(\xi_k)(h_k) + \sum_{t=k+1}^N \alpha^{t-k} \{m(\xi_t)(h_k) - m(\xi_{t-1})(h_k)\} \right]. \tag{S274}$$

For each fixed k , collecting the coefficient of every $m(\xi_t)(h_k)$ gives

$$\begin{aligned}
&m(\xi_k)(h_k) + \sum_{t=k+1}^N \alpha^{t-k} \{m(\xi_t)(h_k) - m(\xi_{t-1})(h_k)\} \\
&= (1-\alpha) \sum_{t=k}^{N-1} \alpha^{t-k} m(\xi_t)(h_k) + \alpha^{N-k} m(\xi_N)(h_k).
\end{aligned} \tag{S275}$$

Indeed, the coefficients are $1-\alpha$ at $t=k$, $\alpha^{t-k} - \alpha^{t+1-k} = (1-\alpha)\alpha^{t-k}$ for $k < t < N$, and α^{N-k} at $t=N$. Interchanging the finite sums once more yields

$$\nabla_{\theta} \hat{\ell}^{\alpha}(\theta) = (1-\alpha) \sum_{t=1}^{N-1} m(\xi_t) \left(\sum_{k=1}^t \alpha^{t-k} h_k \right) + m(\xi_N) \left(\sum_{k=1}^N \alpha^{N-k} h_k \right). \tag{S276}$$

The same algebra applied to Equation (S218) gives

$$\nabla_{\theta} \ell^{\alpha}(\theta) = (1 - \alpha) \sum_{t=1}^{N-1} \rho_t \left(\sum_{k=1}^t \alpha^{t-k} h_k \right) + \rho_N \left(\sum_{k=1}^N \alpha^{N-k} h_k \right). \quad (\text{S277})$$

For every coordinate i , Equation (S205) gives

$$\left\| \sum_{k=1}^t \alpha^{t-k} h_{k,i} \right\|_{\infty} \leq G'(\theta) \sum_{k=1}^t \alpha^{t-k}. \quad (\text{S278})$$

We now prove the bias bound directly from the two terms of the conditional score. Equation (S209) gives

$$\begin{aligned} \left\| \partial_{\theta_i} \log w_{n,j}^{F,\theta} \Big|_{\theta=\phi} \right\|_{\infty} &\leq G'(\theta) \sum_{r=0}^{n-1} \alpha^r, \\ \alpha \left\| \partial_{\theta_i} \log w_{n-1,j}^{F,\theta} \Big|_{\theta=\phi} \right\|_{\infty} &\leq G'(\theta) \sum_{r=1}^{n-1} \alpha^r. \end{aligned} \quad (\text{S279})$$

Apply Lemma S11 to the filtered-weight tangent at time n and, using the second equality in Equation (S211), to α times the filtered-weight tangent at time $n-1$. Since $\text{osc}(f) \leq 2\|f\|_{\infty}$, Equation (S279) gives

$$\begin{aligned} &\left| \mathbb{E} \partial_{\theta_i} \hat{\ell}_n^{\alpha}(\theta) - \partial_{\theta_i} \ell_n^{\alpha}(\theta) \right| \\ &\leq \frac{CG'(\theta)}{J} \left\{ \sum_{r=0}^{n-1} \alpha^r + \sum_{r=1}^{n-1} \alpha^r \right\} \\ &\leq \frac{CG'(\theta)}{J} \sum_{r=0}^{n-1} \alpha^r. \end{aligned} \quad (\text{S280})$$

The last inequality only enlarges C . Summing over n gives

$$\begin{aligned} &\left| \mathbb{E} \partial_{\theta_i} \hat{\ell}^{\alpha}(\theta) - \partial_{\theta_i} \ell^{\alpha}(\theta) \right| \\ &\leq \sum_{n=1}^N \left| \mathbb{E} \partial_{\theta_i} \hat{\ell}_n^{\alpha}(\theta) - \partial_{\theta_i} \ell_n^{\alpha}(\theta) \right| \\ &\leq \frac{CG'(\theta)}{J} \sum_{n=1}^N \sum_{r=0}^{n-1} \alpha^r. \end{aligned} \quad (\text{S281})$$

Squaring, summing over $i = 1, \dots, p$, and taking square roots proves Equation (S272). Finally,

$$\sum_{n=1}^N \sum_{r=0}^{n-1} \alpha^r \leq \sum_{n=1}^N n = \frac{N(N+1)}{2} \leq N^2. \quad (\text{S282})$$

Substitution in Equation (S272) proves Equation (S273). $\square$

Lemma S13 (Variance of the DMOP- α score). *Suppose $J \geq 2$, $\theta = \phi$, $\alpha \in [0, 1]$, Assumptions (A2), (A3), and (A5) hold, and multinomial resampling is used. Then*

$$\text{Var} \left(\nabla_{\theta} \hat{\ell}^{\alpha}(\theta) \right) \lesssim p G'(\theta)^2 \min \left\{ N^2, \frac{1}{J} \sum_{n=1}^N \left(\sum_{r=0}^{n-1} \alpha^r \right)^2 \right\}. \quad (\text{S283})$$

Here $\lesssim$ suppresses the particle-system mixing-time factor. In particular,

$$\begin{aligned} \text{Var} \left(\nabla_{\theta} \hat{\ell}^{\alpha}(\theta) \right) &\lesssim \frac{N p G'(\theta)^2}{J} \left(\sum_{r=0}^{N-1} \alpha^r \right)^2 \\ &= \frac{N p G'(\theta)^2}{J} \begin{cases} \left(\frac{1 - \alpha^N}{1 - \alpha} \right)^2, & 0 \leq \alpha < 1, \\ N^2, & \alpha = 1. \end{cases} \end{aligned} \quad (\text{S284})$$

Apart from the explicitly suppressed particle-system mixing-time factor, the multiplicative constants are independent of N, J , and α , and $\text{Var}(Z) = \mathbb{E} \|Z - \mathbb{E}Z\|_2^2$.

Proof. We next use the fact that the weight tangent is part of the differentiable particle. Equation (S209) and Equation (S205) give, for every coordinate i ,

$$\begin{aligned} \left| \partial_{\theta_i} \log w_{n,j}^{F,\theta} \Big|_{\theta=\phi} \right| &\leq \sum_{r=0}^{n-1} \alpha^r \left| \partial_{\theta_i} \log g_{n-r,j}^{A,F,\theta} \Big|_{\theta=\phi} \right| \\ &\leq G'(\theta) \sum_{r=0}^{n-1} \alpha^r, \end{aligned} \quad (\text{S285})$$

$$\begin{aligned} \left| \partial_{\theta_i} \log w_{n,j}^{P,\theta} \Big|_{\theta=\phi} \right| &= \alpha \left| \partial_{\theta_i} \log w_{n-1,j}^{F,\theta} \Big|_{\theta=\phi} \right| \\ &\leq G'(\theta) \sum_{r=1}^{n-1} \alpha^r. \end{aligned} \quad (\text{S286})$$

The two terms in Equation (S211) are empirical averages of the bounded coordinate functions in Equations (S285) and (S286) of the differentiable particle state. We now apply the precise form of the particle result and write out the filtering step that it does not cover directly. Lemma 2.3 of Karjalainen et al. (2023), with $p = 4$, states that for every bounded real function f of a prediction particle,

$$\|\eta_n^J(f) - \eta_n(f)\|_{L^4} \leq \frac{c(4)}{\sqrt{J}} \text{osc}(f). \quad (\text{S287})$$

This applies directly to the prediction-weight tangent.

For completeness, consider the corresponding filtered average

$$Y = \frac{1}{J} \sum_{j=1}^J f(x_{n,j}^F). \quad (\text{S288})$$

Conditionally on η_n^J , multinomial resampling makes the $x_{n,j}^F$ independent with common law $\pi_n^J = \Psi_n(\eta_n^J)$. If a filtered tangent coordinate is obtained by a deterministic post-selection update, f denotes its composition with that update as a function of the selected prediction particle. Put

$$Z_j = f(x_{n,j}^F) - \pi_n^J(f). \quad (\text{S289})$$

Then, conditionally on η_n^J , the Z_j are independent, $\mathbb{E}(Z_j \mid \eta_n^J) = 0$, and $|Z_j| \leq \text{osc}(f)$. Expanding the fourth power and using conditional independence gives

$$\begin{aligned} & \mathbb{E} \left[|Y - \pi_n^J(f)|^4 \mid \eta_n^J \right] \\ &= \frac{1}{J^4} \left\{ \sum_{j=1}^J \mathbb{E}(Z_j^4 \mid \eta_n^J) + 6 \sum_{1 \leq j < k \leq J} \mathbb{E}(Z_j^2 \mid \eta_n^J) \mathbb{E}(Z_k^2 \mid \eta_n^J) \right\} \\ &\leq \frac{J + 3J(J-1)}{J^4} \text{osc}(f)^4 \\ &\leq \frac{3}{J^2} \text{osc}(f)^4. \end{aligned} \quad (\text{S290})$$

Taking expectations and fourth roots yields

$$\|Y - \pi_n^J(f)\|_{L^4} \leq \frac{3^{1/4}}{\sqrt{J}} \text{osc}(f). \quad (\text{S291})$$

It remains to compare the random weighted prediction measure with the deterministic filter. Since $\pi_n = \Psi_n(\eta_n)$,

$$\begin{aligned} \pi_n^J(f) - \pi_n(f) &= \frac{\eta_n^J(G_n f)}{\eta_n^J(G_n)} - \pi_n(f) \\ &= \frac{\eta_n^J[G_n\{f - \pi_n(f)\}]}{\eta_n^J(G_n)} \\ &= \frac{(\eta_n^J - \eta_n)[G_n\{f - \pi_n(f)\}]}{\eta_n^J(G_n)}, \end{aligned} \quad (\text{S292})$$

where the last equality uses

$$\eta_n[G_n\{f - \pi_n(f)\}] = \eta_n(G_n f) - \pi_n(f)\eta_n(G_n) = 0. \quad (\text{S293})$$

Assumption (A3) gives

$$\begin{aligned} \eta_n^J(G_n) &\geq \underline{G}, \\ \text{osc}[G_n\{f - \pi_n(f)\}] &\leq 2\bar{G} \text{osc}(f). \end{aligned} \quad (\text{S294})$$

Applying Equation (S287) to the centered function in Equation (S292) therefore gives

$$\|\pi_n^J(f) - \pi_n(f)\|_{L^4} \leq \frac{2c(4)\bar{G}}{\underline{G}\sqrt{J}} \text{osc}(f). \quad (\text{S295})$$

Minkowski's inequality and Equations (S291)–(S295) now give the filtered analogue of Karjalainen et al.'s predictor bound:

$$\|Y - \pi_n(f)\|_{L^4} \leq \frac{C}{\sqrt{J}} \text{osc}(f). \quad (\text{S296})$$

For either the prediction or filtered average, let y denote the corresponding deterministic target. Jensen's inequality and the triangle inequality give the required centered form:

$$\begin{aligned} \|Y - \mathbb{E}Y\|_{L^4} &\leq \|Y - y\|_{L^4} + |\mathbb{E}(Y - y)| \\ &\leq 2\|Y - y\|_{L^4}. \end{aligned} \quad (\text{S297})$$

Applying Equations (S287) and (S296) to the two tangent functions, and using $\text{osc}(f) \leq 2\|f\|_\infty$, gives

$$\left\| \frac{1}{J} \sum_{j=1}^J \partial_{\theta_i} \log w_{n,j}^{F,\theta} \Big|_{\theta=\phi} - \mathbb{E} \left[\frac{1}{J} \sum_{j=1}^J \partial_{\theta_i} \log w_{n,j}^{F,\theta} \Big|_{\theta=\phi} \right] \right\|_{L^4} \leq \frac{CG'(\theta)}{\sqrt{J}} \sum_{r=0}^{n-1} \alpha^r, \quad (\text{S298})$$

$$\left\| \frac{1}{J} \sum_{j=1}^J \partial_{\theta_i} \log w_{n,j}^{P,\theta} \Big|_{\theta=\phi} - \mathbb{E} \left[\frac{1}{J} \sum_{j=1}^J \partial_{\theta_i} \log w_{n,j}^{P,\theta} \Big|_{\theta=\phi} \right] \right\|_{L^4} \leq \frac{CG'(\theta)}{\sqrt{J}} \sum_{r=1}^{n-1} \alpha^r. \quad (\text{S299})$$

Minkowski's inequality in L^4 , Equations (S211), (S298), and (S299) imply

$$\begin{aligned} &\left\| \partial_{\theta_i} \hat{\ell}_n^\alpha(\theta) - \mathbb{E} \partial_{\theta_i} \hat{\ell}_n^\alpha(\theta) \right\|_{L^4} \\ &\leq \frac{CG'(\theta)}{\sqrt{J}} \left(\sum_{r=0}^{n-1} \alpha^r + \sum_{r=1}^{n-1} \alpha^r \right) \\ &\leq \frac{CG'(\theta)}{\sqrt{J}} \sum_{r=0}^{n-1} \alpha^r. \end{aligned} \quad (\text{S300})$$

The conditional score at time n is a measurable function of the differentiable particle clouds at times $n-1$ and n . Theorem 3.1 of Karjalainen et al. (2023) applies to this predictor-cloud chain. We spell out the passage from its Dobrushin bound to the score mixing coefficient. Let A be an event generated by the score history through time m , and let B be an event generated by scores from time $m+l$ onward. Let $\mathcal{F}_{m+1}^P$ be the sigma-field generated by all variables through the production of the prediction cloud $X_{m+1,1:J}^P$. The time- m resampling indices enter both the time- m score and the construction of $X_{m+1,1:J}^P$, so $A \in \mathcal{F}_{m+1}^P$. Conditional on this sigma-field, the Markov property gives, for a suitable future-tail kernel K ,

$$\Pr(B \mid \mathcal{F}_{m+1}^P) = K(X_{m+1,1:J}^P, B). \quad (\text{S301})$$

Consequently,

$$\begin{aligned} &|\Pr(A \cap B) - \Pr(A) \Pr(B)| \\ &= |\mathbb{E} [1_A \{K(X_{m+1,1:J}^P, B) - \mathbb{E} K(X_{m+1,1:J}^P, B)\}]| \\ &\leq \sup_{\mathbf{x}, \mathbf{x}'} |K(\mathbf{x}, B) - K(\mathbf{x}', B)| \end{aligned}$$

$$\leq \beta_{\text{TV}}(\mathbf{M}_{m+1,m+l}). \quad (\text{S302})$$

The kernel from the prediction cloud at time $m+1$ to the prediction cloud at time $m+l$ has $l-1$ particle-filter steps; postcomposition with terminal resampling and the future-tail kernel cannot increase its Dobrushin coefficient. Passing to a measurable functional can only reduce the generated σ -algebra. Therefore, for some $\epsilon \in (0,1)$ and $c < \infty$ depending only on the bounds in Assumptions (A2)–(A3), Theorem 3.1 gives, for $l \geq 2$,

$$a(l) \leq (1-\epsilon)^{\lfloor (l-1)/[c\{1+\log J\}] \rfloor}. \quad (\text{S303})$$

For $l=1$, the same display follows from $a(1) \leq 1$. For $m < n$, Davydov's inequality, applied to the centered variables in Equation (S300), gives

$$\begin{aligned} & \left| \text{Cov} \left(\partial_{\theta_i} \hat{\ell}_m^\alpha(\theta), \partial_{\theta_i} \hat{\ell}_n^\alpha(\theta) \right) \right| \\ & \leq C a(n-m)^{1/2} \left\| \partial_{\theta_i} \hat{\ell}_m^\alpha(\theta) - \mathbb{E} \partial_{\theta_i} \hat{\ell}_m^\alpha(\theta) \right\|_{L^4} \left\| \partial_{\theta_i} \hat{\ell}_n^\alpha(\theta) - \mathbb{E} \partial_{\theta_i} \hat{\ell}_n^\alpha(\theta) \right\|_{L^4} \\ & \leq \frac{C G'(\theta)^2}{J} a(n-m)^{1/2} \left(\sum_{r=0}^{m-1} \alpha^r \right) \left(\sum_{r=0}^{n-1} \alpha^r \right) \\ & \leq \frac{C G'(\theta)^2}{2J} a(n-m)^{1/2} \left\{ \left(\sum_{r=0}^{m-1} \alpha^r \right)^2 + \left(\sum_{r=0}^{n-1} \alpha^r \right)^2 \right\}, \end{aligned} \quad (\text{S304})$$

where the last line uses $2xy \leq x^2 + y^2$. Equation (S300) also gives

$$\text{Var} \left(\partial_{\theta_i} \hat{\ell}_n^\alpha(\theta) \right) \leq \frac{C G'(\theta)^2}{J} \left(\sum_{r=0}^{n-1} \alpha^r \right)^2. \quad (\text{S305})$$

Put $q = (1-\epsilon)^{1/2} < 1$. Directly summing the blocked geometric series in Equation (S303) gives

$$\begin{aligned} \sum_{l=1}^{N-1} a(l)^{1/2} & \leq \sum_{l=1}^{N-1} q^{\lfloor (l-1)/[c\{1+\log J\}] \rfloor} \\ & \leq \min \left\{ N-1, \lceil c\{1+\log J\} \rceil \sum_{u=0}^{\infty} q^u \right\} \\ & \leq C \{1 + \min(N, \log J)\}. \end{aligned} \quad (\text{S306})$$

Expanding the variance of the sum, using Equations (S304)–(S306), and observing that for each lag l every squared term occurs at most twice, we obtain

$$\begin{aligned} & \text{Var} \left(\partial_{\theta_i} \hat{\ell}^\alpha(\theta) \right) \\ & = \sum_{n=1}^N \text{Var} \left(\partial_{\theta_i} \hat{\ell}_n^\alpha(\theta) \right) + 2 \sum_{1 \leq m < n \leq N} \text{Cov} \left(\partial_{\theta_i} \hat{\ell}_m^\alpha(\theta), \partial_{\theta_i} \hat{\ell}_n^\alpha(\theta) \right) \\ & \leq \frac{C G'(\theta)^2}{J} \sum_{n=1}^N \left(\sum_{r=0}^{n-1} \alpha^r \right)^2 + \frac{C G'(\theta)^2}{J} \sum_{l=1}^{N-1} a(l)^{1/2} \sum_{n=1}^N \left(\sum_{r=0}^{n-1} \alpha^r \right)^2 \end{aligned}$$

$$\leq \frac{CG'(\theta)^2}{J} \{1 + \min(N, 1 + \log J)\} \sum_{n=1}^N \left(\sum_{r=0}^{n-1} \alpha^r \right)^2. \quad (\text{S307})$$

Summing Equation (S307) over the p coordinates and then using the stated convention for $\lesssim$ gives the stochastic term in Equation (S283). Moreover,

$$\begin{aligned} \sum_{n=1}^N \left(\sum_{r=0}^{n-1} \alpha^r \right)^2 &\leq N \left(\sum_{r=0}^{N-1} \alpha^r \right)^2 \\ &= N \begin{cases} \left(\frac{1 - \alpha^N}{1 - \alpha} \right)^2, & 0 \leq \alpha < 1, \\ N^2, & \alpha = 1, \end{cases} \end{aligned} \quad (\text{S308})$$

Equations (S307) and (S308) prove Equation (S284).

It remains to obtain the other term in the minimum. Taking absolute values in Equation (S276) and using Equation (S278) give, pathwise,

$$\begin{aligned} \left| \partial_{\theta_i} \hat{\ell}^\alpha(\theta) \right| &\leq G'(\theta) \left[\sum_{t=1}^{N-1} (1 - \alpha^t) + \sum_{k=1}^N \alpha^{N-k} \right] \\ &= G'(\theta) \left[(N-1) - \sum_{t=1}^{N-1} \alpha^t + 1 + \sum_{t=1}^{N-1} \alpha^t \right] \\ &= NG'(\theta). \end{aligned} \quad (\text{S309})$$

Consequently,

$$\text{Var} \left(\partial_{\theta_i} \hat{\ell}^\alpha(\theta) \right) \leq \mathbb{E} \left| \partial_{\theta_i} \hat{\ell}^\alpha(\theta) \right|^2 \leq N^2 G'(\theta)^2. \quad (\text{S310})$$

Combining this deterministic bound with Equation (S307), summing over coordinates, and applying the $\lesssim$ convention proves Equation (S283). $\square$

Remark 1 (DMOP-0 and the fixed-lag calculation). *The original DMOP-0 and fixed-lag arguments are given explicitly before Theorem S5. They use the manuscript's convention that $\lesssim$ suppresses the particle-system mixing-time factor. Setting $\alpha = 0$ in Lemma S13 gives*

$$\text{Var} \left(\nabla_{\theta} \hat{\ell}^0(\theta) \right) \lesssim \frac{NpG'(\theta)^2}{J}, \quad (\text{S311})$$

which agrees with that calculation. For fixed $\alpha < 1$, the same lemma gives $N/[J(1 - \alpha)^2]$ up to the common factor $pG'(\theta)^2$. Because the differentiable particle includes its location and tangents, no separate augmented-state or tangent-stability assumption is used. No stronger factor in $1 - \alpha$ is claimed here.

Lemma S14 (Population discounting bias). *Under the conditions of Lemma S10, suppose $\theta = \phi$ and $\alpha \in [0, 1]$. Then*

$$\|\nabla_{\theta} \ell_n^\alpha(\theta) - \nabla_{\theta} \ell_n(\theta)\|_2 \leq C\sqrt{p}G'(\theta) \frac{\varrho(1 - \alpha)}{(1 - \varrho)(1 - \alpha\varrho)}. \quad (\text{S312})$$

Proof. Equation (S218) at α and at $\alpha = 1$ gives

$$\begin{aligned}\nabla_{\theta}\ell_n(\theta) - \nabla_{\theta}\ell_n^{\alpha}(\theta) &= \sum_{r=0}^{n-1} (1 - \alpha^r) \Delta_{n,r} \\ &= \sum_{r=1}^{n-1} (1 - \alpha^r) \Delta_{n,r},\end{aligned}\tag{S313}$$

because $1 - \alpha^0 = 0$. Taking norms, applying the triangle inequality and Lemma S10, and then extending the nonnegative scalar sum to infinity give

$$\begin{aligned}\|\nabla_{\theta}\ell_n(\theta) - \nabla_{\theta}\ell_n^{\alpha}(\theta)\|_2 &\leq \sum_{r=1}^{n-1} (1 - \alpha^r) \|\Delta_{n,r}\|_2 \\ &\leq C\sqrt{p} G'(\theta) \sum_{r=1}^{n-1} (1 - \alpha^r) \varrho^r \\ &\leq C\sqrt{p} G'(\theta) \sum_{r=1}^{\infty} (1 - \alpha^r) \varrho^r \\ &= C\sqrt{p} G'(\theta) \left\{ \sum_{r=1}^{\infty} \varrho^r - \sum_{r=1}^{\infty} (\alpha\varrho)^r \right\} \\ &= C\sqrt{p} G'(\theta) \left\{ \frac{\varrho}{1 - \varrho} - \frac{\alpha\varrho}{1 - \alpha\varrho} \right\} \\ &= C\sqrt{p} G'(\theta) \frac{\varrho(1 - \alpha)}{(1 - \varrho)(1 - \alpha\varrho)}.\end{aligned}\tag{S314}$$

□

Theorem S6 (Refined MSE bound for DMOP- α). *Suppose the observations are fixed, $J \geq 2$, $\theta = \phi$, $\alpha \in [0, 1]$, Assumptions (A2), (A3), and (A5) hold, and multinomial resampling is used. Then there are $C < \infty$ and $\varrho \in (0, 1)$, independent of N, J , and α , such that*

$$\begin{aligned}\mathbb{E} \left\| \nabla_{\theta} \hat{\ell}^{\alpha}(\theta) - \nabla_{\theta} \ell(\theta) \right\|_2^2 \\ \lesssim p G'(\theta)^2 \left[\frac{1}{J} \sum_{n=1}^N \left(\sum_{r=0}^{n-1} \alpha^r \right)^2 + N^2 \left\{ \frac{\varrho(1 - \alpha)}{(1 - \varrho)(1 - \alpha\varrho)} \right\}^2 + \frac{1}{J^2} \left(\sum_{n=1}^N \sum_{r=0}^{n-1} \alpha^r \right)^2 \right].\end{aligned}\tag{S315}$$

Consequently, for $0 \leq \alpha < 1$,

$$\begin{aligned}\mathbb{E} \left\| \nabla_{\theta} \hat{\ell}^{\alpha}(\theta) - \nabla_{\theta} \ell(\theta) \right\|_2^2 \\ \lesssim p G'(\theta)^2 \left\{ \frac{N}{J(1 - \alpha)^2} + N^2(1 - \alpha)^2 + \frac{N^2}{J^2(1 - \alpha)^2} \right\}.\end{aligned}\tag{S316}$$

The finite-sum display in Equation (S315) is defined directly at $\alpha = 1$; its three terms then have orders N^3/J , 0, and N^4/J^2 , respectively.

Proof. For a square-integrable random vector Z and a deterministic vector z , expansion of the square gives

$$\begin{aligned}\mathbb{E}\|Z - z\|_2^2 &= \mathbb{E}\|Z - \mathbb{E}Z + \mathbb{E}Z - z\|_2^2 \\ &= \mathbb{E}\|Z - \mathbb{E}Z\|_2^2 + 2\mathbb{E}\{(Z - \mathbb{E}Z)^\top(\mathbb{E}Z - z)\} + \|\mathbb{E}Z - z\|_2^2 \\ &= \text{Var}(Z) + \|\mathbb{E}Z - z\|_2^2,\end{aligned}\tag{S317}$$

because $\mathbb{E}(Z - \mathbb{E}Z) = 0$.

Apply Equation (S317) with $Z = \nabla_\theta \hat{\ell}^\alpha(\theta)$ and $z = \nabla_\theta \ell(\theta)$. Lemma S13 gives

$$\text{Var}\left(\nabla_\theta \hat{\ell}^\alpha(\theta)\right) \lesssim \frac{pG'(\theta)^2}{J} \sum_{n=1}^N \left(\sum_{r=0}^{n-1} \alpha^r\right)^2.\tag{S318}$$

For the bias, add and subtract the full population DMOP- α score. The triangle inequality, Lemma S12, and Lemma S14 give

$$\begin{aligned}&\left\|\mathbb{E}\nabla_\theta \hat{\ell}^\alpha(\theta) - \nabla_\theta \ell(\theta)\right\|_2 \\ &\leq \left\|\mathbb{E}\nabla_\theta \hat{\ell}^\alpha(\theta) - \nabla_\theta \ell^\alpha(\theta)\right\|_2 \\ &\quad + \left\|\sum_{n=1}^N \{\nabla_\theta \ell_n^\alpha(\theta) - \nabla_\theta \ell_n(\theta)\}\right\|_2 \\ &\leq \frac{C\sqrt{p}G'(\theta)}{J} \sum_{n=1}^N \sum_{r=0}^{n-1} \alpha^r + CN\sqrt{p}G'(\theta) \frac{\varrho(1-\alpha)}{(1-\varrho)(1-\alpha\varrho)}.\end{aligned}\tag{S319}$$

Using $(a+b)^2 \leq 2a^2 + 2b^2$ in Equation (S319) gives

$$\begin{aligned}&\left\|\mathbb{E}\nabla_\theta \hat{\ell}^\alpha(\theta) - \nabla_\theta \ell(\theta)\right\|_2^2 \\ &\leq CpG'(\theta)^2 \left[\frac{1}{J^2} \left(\sum_{n=1}^N \sum_{r=0}^{n-1} \alpha^r\right)^2 + N^2 \left\{ \frac{\varrho(1-\alpha)}{(1-\varrho)(1-\alpha\varrho)} \right\}^2 \right].\end{aligned}\tag{S320}$$

Substitution of Equations (S318) and (S320) into Equation (S317) proves Equation (S315). For $\alpha < 1$,

$$\begin{aligned}\sum_{r=0}^{n-1} \alpha^r &\leq \frac{1}{1-\alpha}, \\ \sum_{n=1}^N \left(\sum_{r=0}^{n-1} \alpha^r\right)^2 &\leq \frac{N}{(1-\alpha)^2}, \\ \sum_{n=1}^N \sum_{r=0}^{n-1} \alpha^r &\leq \frac{N}{1-\alpha}, \\ \frac{\varrho(1-\alpha)}{(1-\varrho)(1-\alpha\varrho)} &\leq \frac{\varrho}{(1-\varrho)^2}(1-\alpha).\end{aligned}\tag{S321}$$

Substitution of Equation (S321) into Equation (S315), with the fixed factors in ϱ absorbed by $\lesssim$, proves Equation (S316). $\square$

Remark 2 (Choice of α within the proved MSE envelope). *The finite-sum bound in Equation (S315) can be minimized directly over $\alpha \in [0, 1]$ for given N, J , and ϱ . There is no general closed form for that finite-sum minimizer. The simplified display gives its rate in the regime where the geometric sums have not saturated at the horizon. Omitting the fixed factors in ϱ , the implicit constants, and the particle-system mixing-time factor suppressed by $\lesssim$, its α -dependent part is*

$$\left(\frac{N}{J} + \frac{N^2}{J^2}\right) (1 - \alpha)^{-2} + N^2(1 - \alpha)^2. \quad (\text{S322})$$

Differentiating Equation (S322) with respect to $1 - \alpha$ gives the rate-balancing choice

$$1 - \alpha \asymp \left(\frac{1}{NJ} + \frac{1}{J^2}\right)^{1/4}. \quad (\text{S323})$$

At this choice, Equation (S322) is of order

$$\frac{N^{3/2}}{\sqrt{J}} \left(1 + \frac{N}{J}\right)^{1/2}. \quad (\text{S324})$$

Thus, up to the common factor $pG'(\theta)^2$ and the suppressed factors, the rate-balancing choice and resulting envelope are

$$1 - \alpha \asymp \begin{cases} J^{-1/2}, & J \lesssim N, \\ (NJ)^{-1/4}, & N \lesssim J \lesssim N^3, \end{cases} \quad \text{MSE} \lesssim pG'(\theta)^2 \begin{cases} N^2/J, & J \lesssim N, \\ N^{3/2}/\sqrt{J}, & N \lesssim J \lesssim N^3. \end{cases} \quad (\text{S325})$$

For $J \gg N^3$, the geometric upper bounds used in Equation (S316) are no longer sharp at the minimizing value. Indeed, when $N(1 - \alpha)$ is small,

$$\sum_{r=0}^{n-1} \alpha^r = n - \frac{n(n-1)}{2}(1 - \alpha) + O\{n^3(1 - \alpha)^2\}. \quad (\text{S326})$$

Substitution into the exact finite-sum envelope shows that its first-order decrease away from $\alpha = 1$ is of order $\{N^4/J + N^5/J^2\}(1 - \alpha)$, whereas the population-discount term is of order $N^2(1 - \alpha)^2$. For $J \gg N^3$, balancing these terms gives

$$1 - \alpha \asymp \frac{N^2}{J}, \quad \text{MSE} \lesssim \frac{pG'(\theta)^2 N^3}{J}. \quad (\text{S327})$$

Combining the regimes gives the rate summary

$$1 - \alpha \asymp \begin{cases} J^{-1/2}, & J \lesssim N, \\ (NJ)^{-1/4}, & N \lesssim J \lesssim N^3, \\ N^2/J, & J \gtrsim N^3, \end{cases} \quad \text{MSE} \lesssim pG'(\theta)^2 \begin{cases} N^2/J, & J \lesssim N, \\ N^{3/2}/\sqrt{J}, & N \lesssim J \lesssim N^3, \\ N^3/J, & J \gtrsim N^3. \end{cases} \quad (\text{S328})$$

The exact envelope also gives a ridge-type endpoint calculation. For $N > 1$, the left derivative at $\alpha = 1$ is

$$\frac{1}{J} \sum_{n=1}^N n^2(n-1) + \frac{N^2(N+1)^2(N-1)}{6J^2} > 0, \quad (\text{S329})$$

so moving left from $\alpha = 1$ decreases the displayed envelope. Its right derivative at $\alpha = 0$ is

$$2 \left\{ \frac{N-1}{J} + \frac{N(N-1)}{J^2} - \frac{N^2 \varrho^2}{1-\varrho} \right\}. \quad (\text{S330})$$

Consequently, if

$$\frac{N-1}{J} + \frac{N(N-1)}{J^2} < \frac{N^2 \varrho^2}{1-\varrho}, \quad (\text{S331})$$

neither endpoint minimizes the displayed envelope, and at least one minimizer lies in $(0, 1)$. These calculations optimize only the proved upper envelope. The hidden constants, suppressed mixing-time factor, and one-sided inequalities prevent them from identifying the minimizer of the actual MSE or proving that the actual MSE is smaller than at both endpoints.

Remark 3 (Equivalent direct L^2 argument). The direct argument does not give a different rate. Minkowski's inequality gives

$$\begin{aligned} & \left\{ \mathbb{E} \left\| \nabla_{\theta} \hat{\ell}^{\alpha}(\theta) - \nabla_{\theta} \ell(\theta) \right\|_2^2 \right\}^{1/2} \\ & \leq \left\{ \text{Var} \left(\nabla_{\theta} \hat{\ell}^{\alpha}(\theta) \right) \right\}^{1/2} + \left\| \mathbb{E} \nabla_{\theta} \hat{\ell}^{\alpha}(\theta) - \nabla_{\theta} \ell^{\alpha}(\theta) \right\|_2 + \left\| \nabla_{\theta} \ell^{\alpha}(\theta) - \nabla_{\theta} \ell(\theta) \right\|_2. \end{aligned} \quad (\text{S332})$$

Substituting Lemmas S13, S12, and S14, then using $(a + b + c)^2 \leq 3(a^2 + b^2 + c^2)$, recovers Equation (S315). Thus no separate direct-MSE proposition is needed.

Remark 4 (Endpoint rates and oscillation normalization). Lemma S13 gives

$$\text{Var} \left(\nabla_{\theta} \hat{\ell}^{\alpha}(\theta) \right) \lesssim \frac{pG'(\theta)^2}{J} \sum_{n=1}^N \left(\sum_{r=0}^{n-1} \alpha^r \right)^2. \quad (\text{S333})$$

Thus the variance orders are N/J at $\alpha = 0$, $N/[J(1-\alpha)^2]$ for fixed $\alpha < 1$, and

$$\frac{1}{J} \sum_{n=1}^N n^2 = \frac{N(N+1)(2N+1)}{6J} = O\left(\frac{N^3}{J}\right) \quad (\text{S334})$$

at $\alpha = 1$, up to the common factor $pG'(\theta)^2$ and the mixing-time factor suppressed by $\lesssim$.

This also resolves the normalization question. A theorem stated for $\text{osc}(f) \leq 1$ must be rescaled when it is applied to the unnormalized score. Since the score has oscillation $O\{NG'(\theta)\}$ at $\alpha = 1$, homogeneity contributes a factor $N^2G'(\theta)^2$ to an MSE bound. If the norm-one path-space theorem has MSE order N/J , as in the sharp historical estimate of Chan et al. (2018), rescaling gives $N^3G'(\theta)^2/J$. If instead one starts from a generic norm-one RMSE bound of order $N/\sqrt{J}$, as in the direct path-space bound of Del Moral and Doucet (2003), rescaling gives

$$\text{RMSE} = O\left(\frac{N^2G'(\theta)}{\sqrt{J}}\right), \quad \text{MSE} = O\left(\frac{N^4G'(\theta)^2}{J}\right). \quad (\text{S335})$$

Therefore an N^4/J MSE can arise from a generic normalized path-functional result, but it is not the sharp current-particle rate proved here.

The lower-bound result of Poyiadjis et al. (2011) concerns their path-space score estimator and does not apply to DMOP-0. At $\alpha = 1$, DMOP has the same undiscounted path-space structure, but its reparameterization score contributions are not asserted to be particlewise identical to the likelihood-ratio score contributions of Poyiadjis et al. (2011). More precisely, in the present notation, Theorem 1 of Poyiadjis et al. (2011) gives a lower bound of order N^2 for the asymptotic variance of $\sqrt{J}$ times the error of their $O(J)$ path-space estimator. Its leading Monte Carlo MSE in the particle central-limit regime is therefore at least of order N^2/J . The generic path-space L^2 estimate quoted there from Del Moral and Doucet (2003) gives the distinct upper bound $O(N^4/J)$. These two results do not identify a single sharp horizon order for the MSE of that estimator; in particular, the N^4/J quantity is an upper bound, not the Poyiadjis lower-bound rate. Their separate $O(J^2)$ marginal algorithm is not the path-space estimator covered by Theorem 1.

Remark 5 (Scope of the proved α dependence). *The supplement proves the $(1 - \alpha)^{-2}$ factor in Equation (S333). It does not claim the stronger $(1 - \alpha)^{-1}$ factor. Because the differentiable particle already contains its location and tangents, this limitation is not a tangent-stability caveat; it is simply the limit of the covariance bound used in Lemma S13.*

Remark 6 (Ridge-type tuning and endpoint derivatives). *At $\theta = \phi$, the particle genealogy in Equation (S214) does not depend on α . Hence*

$$\text{MSE}_\alpha = \mathbb{E} \left\| \sum_{n=1}^N \sum_{r=0}^{n-1} \alpha^r \Delta_{n,r}^J - \nabla_\theta \ell(\theta) \right\|_2^2 \quad (\text{S336})$$

is a polynomial in α and has no singularity at $\alpha = 1$. Differentiating the finite sum gives

$$\frac{d}{d\alpha} \text{MSE}_\alpha = 2\mathbb{E} \left[\left\{ \sum_{n=1}^N \sum_{r=0}^{n-1} \alpha^r \Delta_{n,r}^J - \nabla_\theta \ell(\theta) \right\}^\top \left\{ \sum_{n=1}^N \sum_{r=1}^{n-1} r \alpha^{r-1} \Delta_{n,r}^J \right\} \right]. \quad (\text{S337})$$

Therefore moving right from DMOP-0 lowers the actual MSE precisely when the right derivative in Equation (S337) at zero is negative; moving left from DMOP-1 lowers it precisely when the left derivative at one is positive. If both inequalities are strict, continuity gives an interior minimizer whose MSE is smaller than the MSE at either endpoint. The assumptions used here do not determine these signed covariances, so they do not prove either endpoint sign.

Remark 2 gives the corresponding calculation for the proved upper envelope. That calculation can establish an interior minimizer of the displayed envelope under Equation (S331), but it does not determine either endpoint derivative of the actual MSE in Equation (S337).

Remark 7 (Endpoint interpretation). *There is no DMOP-0 MSE lower bound under the present assumptions: conditional population biases can cancel and can vanish in special models. Consequently, upper bounds alone cannot prove that an interior choice of α improves on DMOP-0.*

For every fixed N , the finite-sum bound at $\alpha = 1$ has the ordinary Monte Carlo variance order J^{-1} ; its dependence on the horizon is N^3/J . The squared particle-bias term is N^4/J^2 and is lower order as $J \rightarrow \infty$ with N fixed. These fixed- N statements should not be confused with joint limits in which N and J grow together.

The quadratic asymptotic-variance lower bound of Poyiadjis et al. (2011) is already stated for an unnormalized additive functional. It must not be multiplied by another N^2 normalization factor. Moreover, the DMOP-1 reparameterization score contributions are not generally particlewise identical to their likelihood-ratio score contributions, so their lower bound is not automatically a DMOP-1 lower bound.

S8 Figures for Bayesian Inference




Figure S-1: Convergence diagnostics for the Metropolis-Hastings variant particle MCMC with the random walk proposal and an informative empirical prior from IF2, on the trend parameter in the Dhaka cholera model of King et al. (2008). The poor convergence shown here, using an MCMC algorithm that does not have access to derivatives, contrasts with Fig. 6 in the main text.

MCMC Diagnostics




Figure S-2: Convergence diagnostics for a No-U-Turn Sampler (NUTS) with uniform priors on a compact set, on the trend parameter in the Dhaka cholera model of King et al. (2008). The NUTS sampler explores more of the posterior than the Metropolis-Hastings sampler, but fails to converge quickly. This poor convergence, compared with Fig. 6 in the main text, demonstrates the value of the empirical prior for obtaining successful convergence.




Figure S-3: Posterior estimates from NUTS powered by MOP- α with the informative nonparametric empirical prior obtained from a kernel density estimator on the IF2 parameter cloud. The posterior estimates from each chain are largely in agreement.

S9 Justification for the Choice of Baseline Comparisons

The reasoning we use in the main text to support the magnitude of our contribution is as follows: (i) iterated filtering methods are the current state of the art, and have been for a decade; (ii) we have demonstrated a substantial advance over iterated filtering. The main text focuses on demonstrating (ii), while supporting (i) via references. Listing a few references amounts only to anecdotal evidence, which is not sufficient to persuade someone properly skeptical about the claim. Therefore, this supplement provides an additional quantitative evaluation.

We count the number of papers published in three top scientific journals (*Science*, *Nature*, and *Proceedings of the National Academy of Sciences of the USA (PNAS)*) that carry out statistical inference for partially observed nonlinear non-Gaussian stochastic dynamic models to address a scientific question by fitting mechanistic models to time series data. For all methods, papers in these high-profile journals are expected to represent only be the tip of the iceberg, but this count provides a concrete measure of the capability of the methods to handle world-class scientific inference problems. We could have considered additional journals, but the three we chose are undoubtedly prestigious and their scope covers the full range of scientific endeavor so there is less room for subfield bias. For iterated filtering, our high-impact application count gives the following: 1. He et al. (2023), 2. Fox et al. (2022), 3. Subramanian et al. (2021), 4. Pei et al. (2021), 5. Li et al. (2020), 6. Giezendanner et al. (2020), 7. Wale et al. (2019), 8. Pons-Salort and Grassly (2018), 9. Ranjeva et al. (2017), 10. Martinez et al. (2016), 11. Becker et al. (2016), 12. Bakker et al. (2016), 13. Laneri et al. (2015), 14. Blake et al. (2014), 15. Blackwood et al. (2013a), 16. Blackwood et al. (2013b), 17. King et al. (2008),

A referee advised consideration of an alternative approach by Gu and Kong (1998) that we found to have a high-impact application count of zero, via Google Scholar.

A referee also recommended comparison with Lindsten et al. (2014), for which the high-impact application count is also zero. Lindsten et al. (2014) has been employed once in *Nature Communications* (Calafat et al., 2018) though that journal did not quite make it onto our elite list. We also found five applications of iterated filtering in *Nature Communications* (Ranjeva et al., 2019; Rodó et al., 2021; Santos-Vega et al., 2022; Briga et al., 2024; Perez-Saez et al., 2025). There is a paper in the Statistics section of *PNAS* by Wang et al. (2022) that demonstrates the method of Lindsten et al. (2014), but here we are counting the demonstrated ability of the methodology to generate high-profile scientific papers, not statistical methodology papers. Otherwise, we would also count Ionides et al. (2006) and Ionides et al. (2015) for iterated filtering.

This reasoning supports iterated filtering as an empirically justified point of reference, in preference to the approaches of Lindsten et al. (2014) and Gu and Kong (1998). Further, the method of Lindsten et al. (2014) does not have the plug-and-play property so it cannot readily be applied to our example. The method of Gu and Kong (1998) would have to be combined with an approach such as Andrieu et al. (2010) that has known scalability problems. Therefore, we focus on comparing with iterated filtering and we do not provide a comparison with Lindsten et al. (2014) or Gu and Kong (1998).

Benchmark tasks are invaluable for making progress on challenging problems (Donoho, 2017). The Dhaka cholera data and model have served as such a benchmark since its introduction by King et al. (2008). Other papers testing methods on this example include Ionides et al. (2015); Fasiolo et al. (2016); Wycoff et al. (2026). Thus, our test is established as a demanding inference problem that challenges the best available methods. We note that this benchmark challenge contains various difficulties that arise commonly in scientific inquiry but are absent from simple SIR models such as

the simulation study considered by Lindsten et al. (2014). Ability to accommodate such features is critical for the methodology to be effective for scientific inquiry.

Iterated filtering has recently been used as a benchmark for new POMP model methodology by Whitehouse et al. (2023). This analysis presented numerical results that claimed to beat iterated filtering, but was subsequently found to contain numerical errors. Corrected results are favorable for iterated filtering methods (Hao et al., 2024; Whitehouse et al., 2025; He et al., 2025).

Recent results of Chen et al. (2025) have extended the previous theoretical understanding of iterated filtering. The ongoing study of iterated filtering as a state-of-the-art approach to inference for POMP models gives additional support for our decision to use iterated filtering as a benchmark for our new methodology.

S10 Algorithmic Parameters

We include the full set of algorithmic parameters used for the Dhaka results in Table S-1. The warm start iterations are IF2 iterations, which are followed by gradient descent (GD) iterations when using IFAD. The IF2 cooling rate is the fractional reduction in the parameter perturbations over 50 iterations, so that a cooling rate of 0.5 corresponds to approximately 1% cooling per iteration. The random walk standard deviation (RWSD) scales the noise added at each filtering step to each regular parameter. Initial value parameters (IVPs), that affect the value of the latent process only at time t_0 , are treated differently: they are perturbed only at time t_0 , with a larger RWSD. Parameters are transformed to a unit scale, for example by taking logarithms or a logistic transformation, so that a single RWSD value works for all parameters.

Hyperparameter	Standard IF2	IFAD-0	IFAD-0.97	IFAD-1
Number of Particles (J)	5,000	5,000	5,000	5,000
Warm-start Iterations	650	60	175	60
GD Iterations	—	225	175	225
IF2 Cooling Rate	0.8	0.5	0.5	0.5
Initial RWSD	0.02	0.02	0.02	0.02
Initial RWSD (IVPs)	0.16	0.16	0.16	0.16
Gradient Optimizer	—	Adam	Adam	Adam
Momentum (β_1)	—	0.0	0.9	0.9
LR Schedule	—	Cosine Decay	Cosine Decay	Cosine Decay
Initial LR (η_0)	—	0.1	0.1	0.1
Warm-up Steps	—	10	10	10
Eval. Particles / Reps	5,000 / 36	5,000 / 36	5,000 / 36	5,000 / 36

Table S-1: Algorithmic parameters for the global search benchmarking on the Dhaka cholera model.

Following some experimentation, IFAD was implemented by plugging the DMOP gradient estimate into the Adam optimizer (Kingma and Ba, 2014). The momentum parameter in Adam was turned off for IFAD-0, since this low-variance estimate did not require it. Decreasing the learning rate (LR) was found to be helpful, and we selected the commonly used cosine decay. Warm-up steps were taken with Adam to stabilize its initial performance and avoid noisy

Supplementary References

- Andrieu, C., Doucet, A., and Holenstein, R. (2010). Particle Markov chain Monte Carlo methods. *Journal of the Royal Statistical Society, Series B*, 72:269–342.
- Bakker, K. M., Martinez-Bakker, M. E., Helm, B., and Stevenson, T. J. (2016). Digital epidemiology reveals global childhood disease seasonality and the effects of immunization. *Proceedings of the National Academy of Sciences*, 113(24):6689–6694.
- Becker, A. D., Birger, R. B., Teillant, A., Gastanaduy, P. A., Wallace, G. S., and Grenfell, B. T. (2016). Estimating enhanced prevaccination measles transmission hotspots in the context of cross-scale dynamics. *Proceedings of the National Academy of Sciences*, 113(51):14595–14600.
- Blackwood, J. C., Cummings, D. A. T., Broutin, H., Iamsirithaworn, S., and Rohani, P. (2013a). Deciphering the impacts of vaccination and immunity on pertussis epidemiology in Thailand. *Proceedings of the National Academy of Sciences*, 110:9595–9600.
- Blackwood, J. C., Streicker, D. G., Altizer, S., and Rohani, P. (2013b). Resolving the roles of immunity, pathogenesis, and immigration for rabies persistence in vampire bats. *Proceedings of the National Academy of Sciences*.
- Blake, I. M., Martin, R., Goel, A., Khetsuriani, N., Everts, J., Wolff, C., Wassilak, S., Aylward, R. B., and Grassly, N. C. (2014). The role of older children and adults in wild poliovirus transmission. *Proceedings of the National Academy of Sciences*, 111(29):10604–10609.
- Briga, M., Goult, E., Brett, T. S., Rohani, P., and Domenech de Cellès, M. (2024). Maternal pertussis immunization and the blunting of routine vaccine effectiveness: A meta-analysis and modeling study. *Nature Communications*, 15(1):921.
- Calafat, F. M., Wahl, T., Lindsten, F., Williams, J., and Frajka-Williams, E. (2018). Coherent modulation of the sea-level annual cycle in the United States by Atlantic Rossby waves. *Nature Communications*, 9(1):2571.
- Chan, H. P., Del Moral, P., and Jasra, A. (2018). A sharp first order analysis of Feynman–Kac particle models. *Stochastic Processes and their Applications*, 128(1):332–353.
- Chen, Y., Gerber, M., Andrieu, C., and Douc, R. (2025). Self-organizing state-space models with artificial dynamics. *Journal of the Royal Statistical Society, Series B*, page qkaf069.
- Chopin, N. (2004). Central limit theorem for sequential Monte Carlo methods and its application to Bayesian inference. *Annals of Statistics*, 32:2385–2411.
- Chopin, N. and Papaspiliopoulos, O. (2020). *An Introduction to Sequential Monte Carlo*. Springer.
- Del Moral, P. (2004). *Feynman-Kac Formulae: Genealogical and Interacting Particle Systems with Applications*. Springer, New York.
- Del Moral, P. and Doucet, A. (2003). On a class of genealogical and interacting Metropolis models. In Azéma, J., Émery, M., Ledoux, M., and Yor, M., editors, *Séminaire de Probabilités XXXVII*, pages 415–446. Springer, Berlin.

- Del Moral, P. and Rio, E. (2011). Concentration inequalities for mean field particle models. *The Annals of Applied Probability*, 21(3).
- Donoho, D. (2017). 50 years of data science. *Journal of Computational and Graphical Statistics*, 26(4):745–766.
- Fasiolo, M., Pya, N., and Wood, S. N. (2016). A comparison of inferential methods for highly nonlinear state space models in ecology and epidemiology. *Statistical Science*, 31(1):96–118.
- Fox, S. J., Lachmann, M., Tec, M., Pasco, R., Woody, S., Du, Z., Wang, X., Ingle, T. A., Javan, E., Dahan, M., Gaither, K., Escott, M. E., Adler, S. I., Johnston, S. C., Scott, J. G., and Meyers, L. A. (2022). Real-time pandemic surveillance using hospital admissions and mobility data. *Proceedings of the National Academy of Sciences*, 119(7):e2111870119.
- Giezendanner, J., Pasetto, D., Perez-Saez, J., Cerrato, C., Viterbi, R., Terzago, S., Palazzi, E., and Rinaldo, A. (2020). Earth and field observations underpin metapopulation dynamics in complex landscapes: Near-term study on carabids. *Proceedings of the National Academy of Sciences*, 117(23):12877–12884.
- Gu, M. G. and Kong, F. H. (1998). A stochastic approximation algorithm with Markov chain Monte-Carlo method for incomplete data estimation problems. *Proceedings of the National Academy of Sciences*, 95(13):7270–7274.
- Hao, Y., Abkemeier, A. J., and Ionides, E. L. (2024). Poisson approximate likelihood compared to the particle filter. *arXiv:2409.12173*.
- He, D., Lin, L., Artzy-Randrup, Y., Demirhan, H., Cowling, B. J., and Stone, L. (2023). Resolving the enigma of Iquitos and Manaus: A modeling analysis of multiple COVID-19 epidemic waves in two Amazonian cities. *Proceedings of the National Academy of Sciences*, 120(10):e2211422120.
- He, K., Hao, Y., and Ionides, E. L. (2025). Poisson Approximate Likelihood versus the block particle filter for a spatiotemporal measles model. *arXiv:2507.09121*.
- Ionides, E. L., Bretó, C., and King, A. A. (2006). Inference for nonlinear dynamical systems. *Proceedings of the National Academy of Sciences*, 103:18438–18443.
- Ionides, E. L., Nguyen, D., Atchadé, Y., Stoev, S., and King, A. A. (2015). Inference for dynamic and latent variable models via iterated, perturbed Bayes maps. *Proceedings of the National Academy of Sciences*, 112(3):719–724.
- Karjalainen, J., Lee, A., Singh, S. S., and Vihola, M. (2023). On the forgetting of particle filters. *arXiv:2309.08517*.
- King, A. A., Ionides, E. L., Pascual, M., and Bouma, M. J. (2008). Inapparent infections and cholera dynamics. *Nature*, 454:877–880.
- Kingma, D. P. and Ba, J. (2014). Adam: A method for stochastic optimization. *arXiv:1412.6980*.
- Laneri, K., Paul, R. E., Tall, A., Faye, J., Diene-Sarr, F., Sokhna, C., Trape, J.-F., and Rodó, X. (2015). Dynamical malaria models reveal how immunity buffers effect of climate variability. *Proceedings of the National Academy of Sciences*, 112(28):8786–8791.

- Li, R., Pei, S., Chen, B., Song, Y., Zhang, T., Yang, W., and Shaman, J. (2020). Substantial undocumented infection facilitates the rapid dissemination of novel coronavirus (SARS-CoV-2). *Science*, 368(6490):489–493.
- Lindsten, F., Jordan, M. I., and Schön, T. B. (2014). Particle Gibbs with ancestor sampling. *The Journal of Machine Learning Research*, 15(1):2145–2184.
- Martinez, P. P., King, A. A., Yunus, M., Faruque, A., and Pascual, M. (2016). Differential and enhanced response to climate forcing in diarrheal disease due to rotavirus across a megacity of the developing world. *Proceedings of the National Academy of Sciences*, 113(15):4092–4097.
- Naesseth, C., Linderman, S., Ranganath, R., and Blei, D. (2018). Variational sequential Monte Carlo. In Storkey, A. and Perez-Cruz, F., editors, *Proceedings of the Twenty-First International Conference on Artificial Intelligence and Statistics*, volume 84 of *Proceedings of Machine Learning Research*, pages 968–977. PMLR.
- Pei, S., Liljeros, F., and Shaman, J. (2021). Identifying asymptomatic spreaders of antimicrobial-resistant pathogens in hospital settings. *Proceedings of the National Academy of Sciences*, 118(37):e2111190118.
- Perez-Saez, J., Bellon, M., Lessler, J., Berthelot, J., Hodcroft, E. B., Michielin, G., Pennacchio, F., Lamour, J., Laubscher, F., L’Huillier, A. G., Posfay-Barbe, K. M., Maerkl, S. J., Guessous, I., Azman, A. S., Eckerle, I., Stringhini, S., and Lortie, E. (2025). Evolving infectious disease dynamics shape school-based intervention effectiveness. *Nature Communications*, 16(1):6597.
- Pons-Salort, M. and Grassly, N. C. (2018). Serotype-specific immunity explains the incidence of diseases caused by human enteroviruses. *Science*, 361(6404):800–803.
- Poyiadjis, G., Doucet, A., and Singh, S. S. (2011). Particle approximations of the score and observed information matrix in state space models with application to parameter estimation. *Biometrika*, 98(1):65–80.
- Ranjeva, S., Subramanian, R., Fang, V. J., Leung, G. M., Ip, D. K., Perera, R. A., Peiris, J. M., Cowling, B. J., and Cobey, S. (2019). Age-specific differences in the dynamics of protective immunity to influenza. *Nature Communications*, 10(1):1660.
- Ranjeva, S. L., Baskerville, E. B., Dukic, V., Villa, L. L., Lazcano-Ponce, E., Giuliano, A. R., Dwyer, G., and Cobey, S. (2017). Recurring infection with ecologically distinct HPV types can explain high prevalence and diversity. *Proceedings of the National Academy of Sciences*, 114(51):13573–13578.
- Rodó, X., Martinez, P. P., Siraj, A., and Pascual, M. (2021). Malaria trends in Ethiopian highlands track the 2000 ‘slowdown’ in global warming. *Nature Communications*, 12(1):1555.
- Roosta-Khorasani, F. and Mahoney, M. W. (2016). Sub-sampled Newton methods I: Globally convergent algorithms. *arXiv:1601.04737*.
- Santos-Vega, M., Martinez, P., Vaishnav, K., Kohli, V., Desai, V., Bouma, M., and Pascual, M. (2022). The neglected role of relative humidity in the interannual variability of urban malaria in Indian cities. *Nature Communications*, 13(1):533.

- Ścibior, A. and Wood, F. (2021). Differentiable particle filtering without modifying the forward pass. *arXiv:2106.10314*.
- Subramanian, R., He, Q., and Pascual, M. (2021). Quantifying asymptomatic infection and transmission of COVID-19 in New York City using observed cases, serology, and testing capacity. *Proceedings of the National Academy of Sciences*, 118(9).
- Wale, N., Jones, M. J., Sim, D. G., Read, A. F., and King, A. A. (2019). The contribution of host cell-directed vs. parasite-directed immunity to the disease and dynamics of malaria infections. *Proceedings of the National Academy of Sciences*, 116(44):22386–22392.
- Wang, E. T., Vannucci, M., Haneef, Z., Moss, R., Rao, V. R., and Chiang, S. (2022). A Bayesian switching linear dynamical system for estimating seizure chronotypes. *Proceedings of the National Academy of Sciences*, 119(46):e2200822119.
- Whitehouse, M., Whiteley, N., and Rimella, L. (2023). Consistent and fast inference in compartmental models of epidemics using Poisson Approximate Likelihoods. *Journal of the Royal Statistical Society, Series B*, 85(4):1173–1203.
- Whitehouse, M., Whiteley, N., and Rimella, L. (2025). Correction to: Consistent and fast inference in compartmental models of epidemics using Poisson approximate likelihoods. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, pre-published online.
- Wycoff, N., Smith, J. W., Booth, A. S., and Gramacy, R. B. (2026). Voronoi candidates for Bayesian optimization. *Journal of Global Optimization*, 94:175–201.